

CSCB09: Software Tools and Systems Programming

Bianca Schroeder

bianca@cs.toronto.edu

IC 460

Who am I

- Professor with research area in computer systems / system reliability
- Lots of research collaborations with companies (e.g. Google, Network Appliances) and supercomputing centers.
- Associate department chair for computer science
- Teaching: B09 or C69, grad courses on reliability and storage systems.

Administrivia - Communications

- Main constraint:
 - There is 200 of you, but only one of me ...
- Therefore:
 - Step 1: Talk to your TA in the tutorial
 - Step 2: Post on Piazza
 - Piazza (you should have received an invitation)
 - Make sure to read their “Terms of Use” and contact me if you are not comfortable
 - Only if step 1 & 2 fail: email me
 - Email: bianca@cs.toronto.edu
 - Email must include your name.
 - Subject must include B09
 - Email is a formal method of communication:
 - State your question clearly, with enough context.
 - **Sign it** (Name and utorid are the most useful.)

Course Information

- Check the course information sheet (on the course web page)
- The course web page and Piazza are the official sources of announcements.

<http://www.cs.toronto.edu/~bianca/cscb09s17/>

<https://piazza.com/class/j29k4eb91xu7jy>

- It's your responsibility to follow Piazza (e.g. find out about changes to assignment deadlines & requirements)

Course work

- A1: Shell and very basic C programming
 - A2: Pointers, memory management (C)
 - A3: Fork and pipes (C)
-
- PCRS work
 - One midterm, one final.
 - Final: Must receive $\geq 40\%$ to pass course

Assignments

- A1: Shell and very basic C programming
- A2: Pointers, memory management (C)
- A3: Fork and pipes (C)

- Assignments:
 - Submission through SVN and Markus
 - All code **must** work on the lab machines (BV473/mathlab) to receive full marks.
 - Code that does not compile on lab machines will receive **0**.
 - Code that does not run will not get any marks from automarking
 - *Don't wait until the last day!*

Did you catch that?

Code that does not compile
will receive a grade of
0

(Almost) Weekly Tutorials

- Labs
 - Starting next week, in BV 473
 - Work in pairs
 - To prepare you for the assignments

The inverted classroom



- Preparation before class (videos+exercises) on PCRS:
<https://cms-pcrs.utsc.utoronto.ca/cscb09/content/quests>
- Hands-on activities in class
 - Need to bring computers to class (or share)

The inverted classroom



- For next week work on “Week 2 preparation” on PCRS:
<https://cms-pcrs.utsc.utoronto.ca/cscb09/content/quests>
- Complete all modules that are marked as “Credit”, i.e.:
 - “Input, Output and Compiling”, “Iteration”, “Arrays”.

Plagiarism

- “The work you submit must be your own, done without participation by others. It is an academic offense to hand in anything written by someone else without acknowledgement.”
- You are not helping your friend when you *give* him or her a copy of your assignment.
- You are hurting your friend when you *ask* him or her to give you a copy of their assignment.

What is cheating?

- Cheating is
 - copying parts or all of another student's assignment
 - including code from books, web sites, other courses without attribution
 - getting someone else to do substantial parts of your assignment
 - giving someone else your solution
- Cheating is not
 - helping to find a bug in a friend's code (be careful)
 - helping each other understand man pages or example code.
- We will run a plagiarism detector on each assignment!

Course Overview

- Unix Software Tools
 - Understanding the shell
 - Shell programming
 - Make
- Systems Programming
 - C
 - files
 - processes
 - concurrency
 - communication

Windows users

- You can install **cygwin**, which gives you a linux-like environment on windows.
- All code must run on the lab machines, so if you develop C code elsewhere make sure it works on the lab machines (or matlab).
- If you want to work remotely on the lab machines or matlab, you can **install putty** on windows, which allows you to connect remotely. On linux machines you can use **ssh**.
- As an instructor I cannot provide tech-support for any of the above. However, you may ask ¹⁴ Tianze in IC464 or your TA.

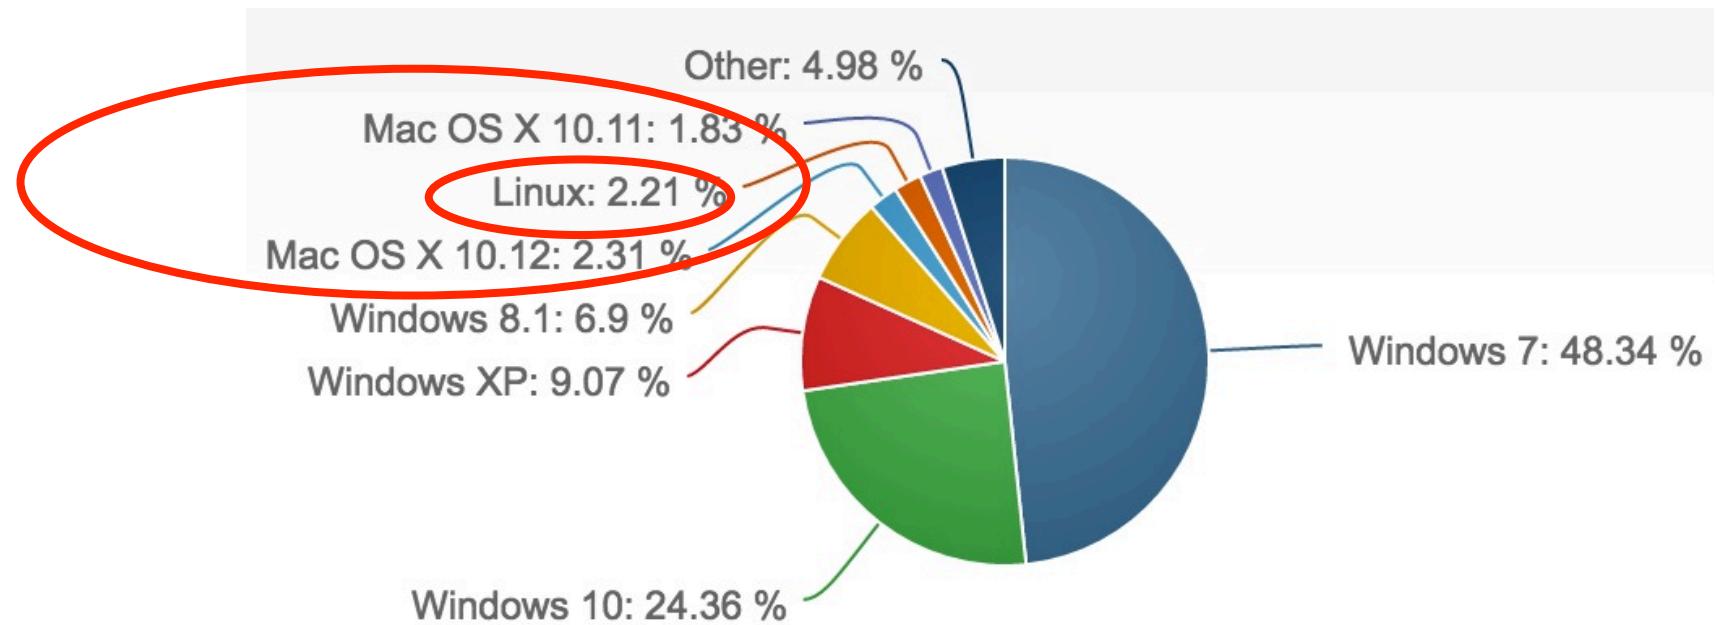
Self Study Topics

- Using SVN
- Using basic Unix
- Learning an editor – vi, emacs, scite, nedit, ...
- We will spend a bit of time on these in class/tutorial, but not a lot.

The plan for today

- What is Unix and why are we using it?
- What is an operating system?
- What are files and directories?
- What is a shell and why do I care about shell programming?
- Along the way: learn/refresh basic Unix commands

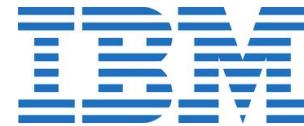
Is anybody using Unix/Linux?



- %age of the world's laptops/desktops running Unix is smallish ...

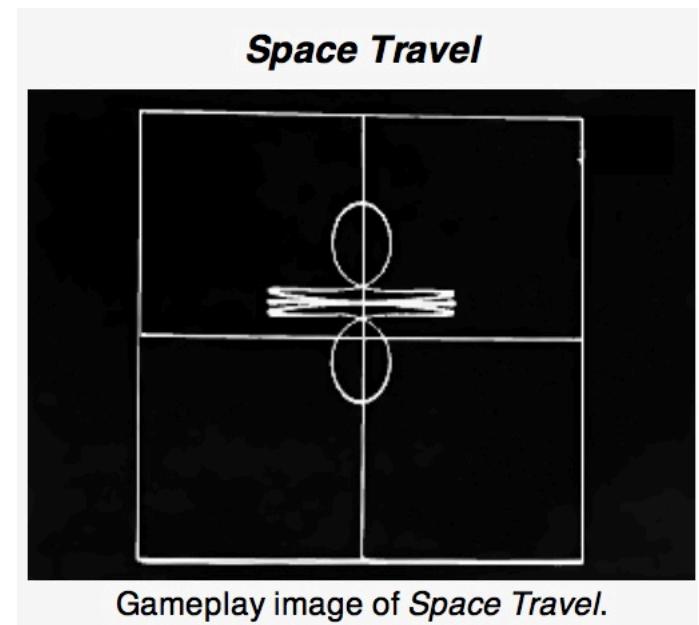
Is anybody using Unix/Linux?

- %age of laptops/desktops running Unix?
 - 8%
- %age of tablets/smartphones running Unix?
 - 95%
- %age of web servers running Unix?
 - 65%
- %age of supercomputers running Unix?
 - 98%



Why Unix?

- 1969:
 - Ken Thompson at Bell Labs wanted to play Space Travel on his DEC PDP-7
 - Wrote the first version in assembler in one month
- Now:
 - Available on many platforms
 - Multi-user, multi-programmed
 - Good at
 - sharing computer resources
 - manipulation of files, processes, programs
 - inter-process and inter-machine communication

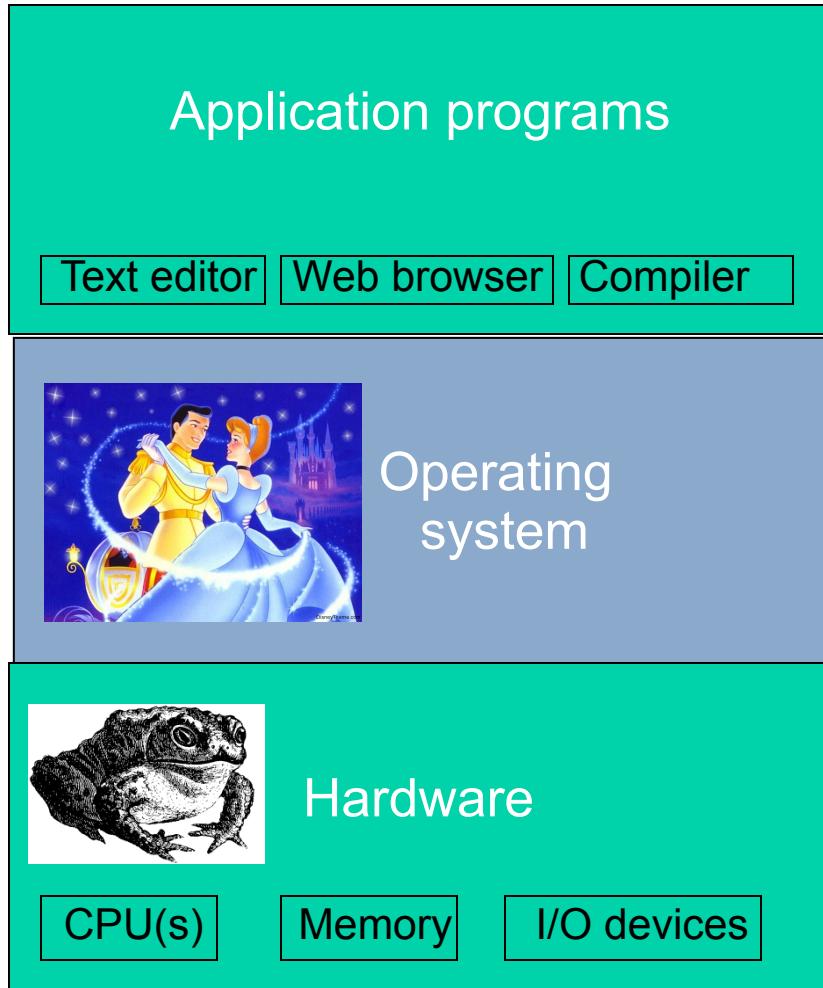


Why Unix?



- Open source
 - Often free
 - Access to operating features
 - Can be customized

What is an operating system?

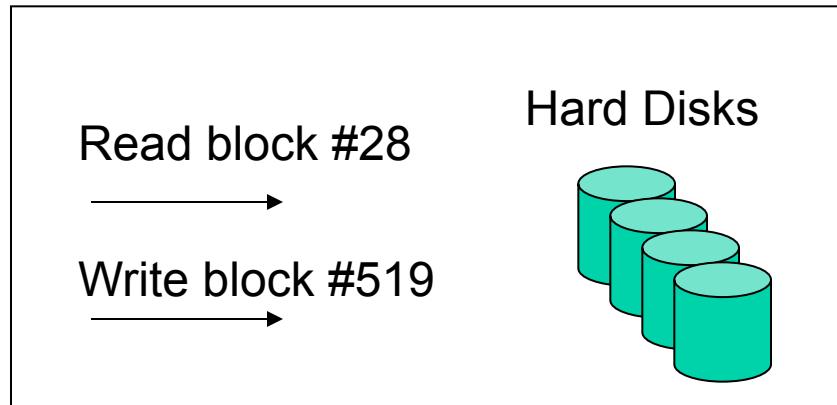


The software layer between user applications and hardware.

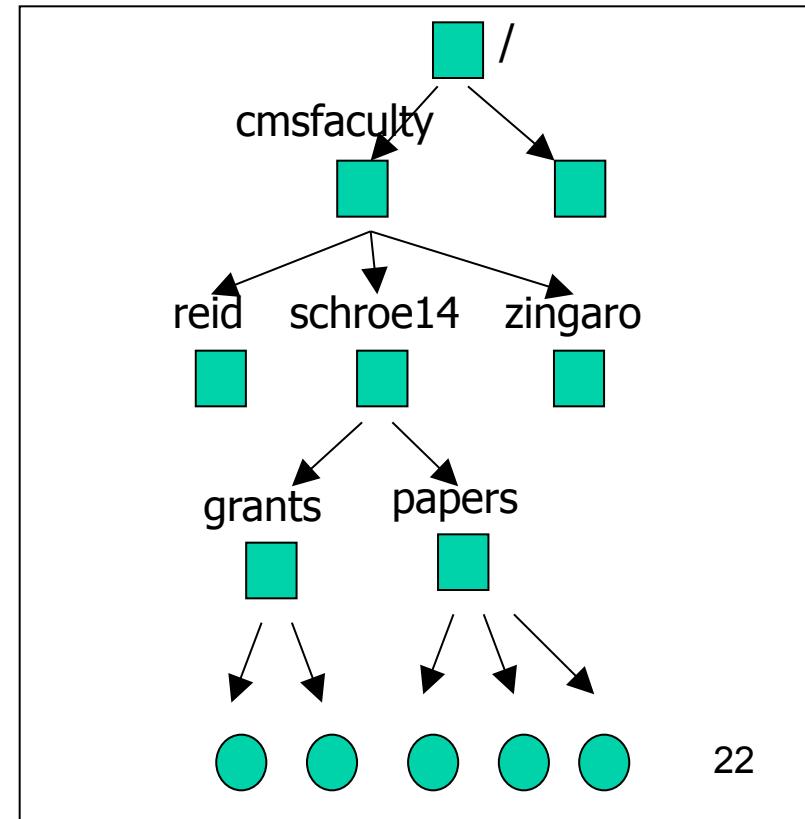
- Serves as a resource manager
 - allows proper use of resources (hardware, software, data)
- Serves as a control program (protection)
 - controls execution of user programs to prevent errors and improper use of the computer
- Turns ugly hardware into beautiful abstractions (provides services)

One such abstraction: Files and Directories

Reality

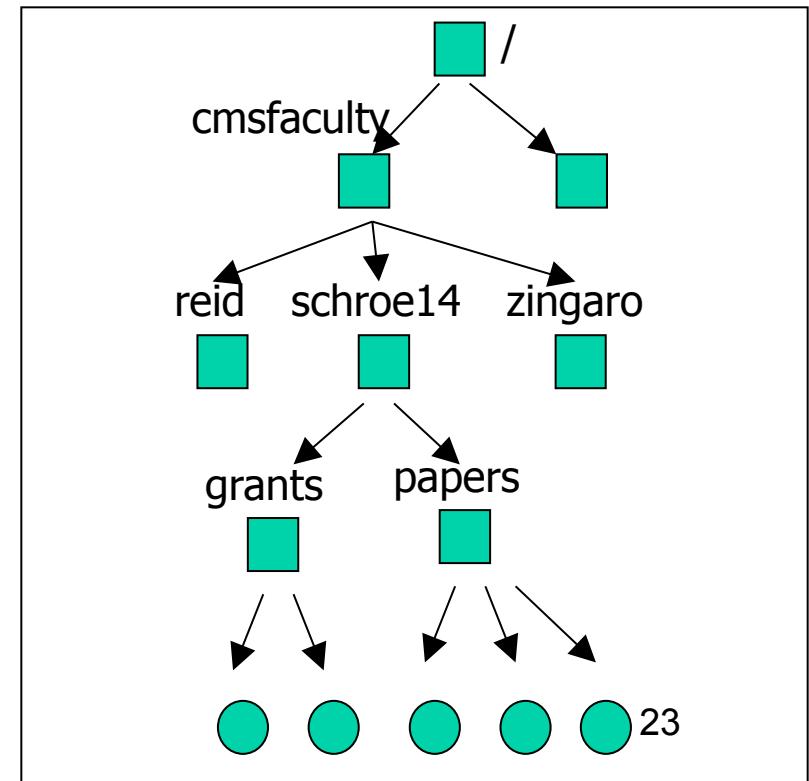


Abstraction



Unix commands to navigate directories and files

- Unix has the notion of a *current working directory*. Get it by typing `pwd` in the shell.
 - Type `ls` to see all files and directories in current working directory.
 - Use `cd` to change your working directory.
 - Use `cat` or `more` to view file contents.
 - Use `cp`, `mv`, `rm` to copy / move/ remove files.
 - Use `mkdir` to create a directory.
- (Use man pages to learn more about commands, e.g. type “`man ls`” for info on the `ls` command)

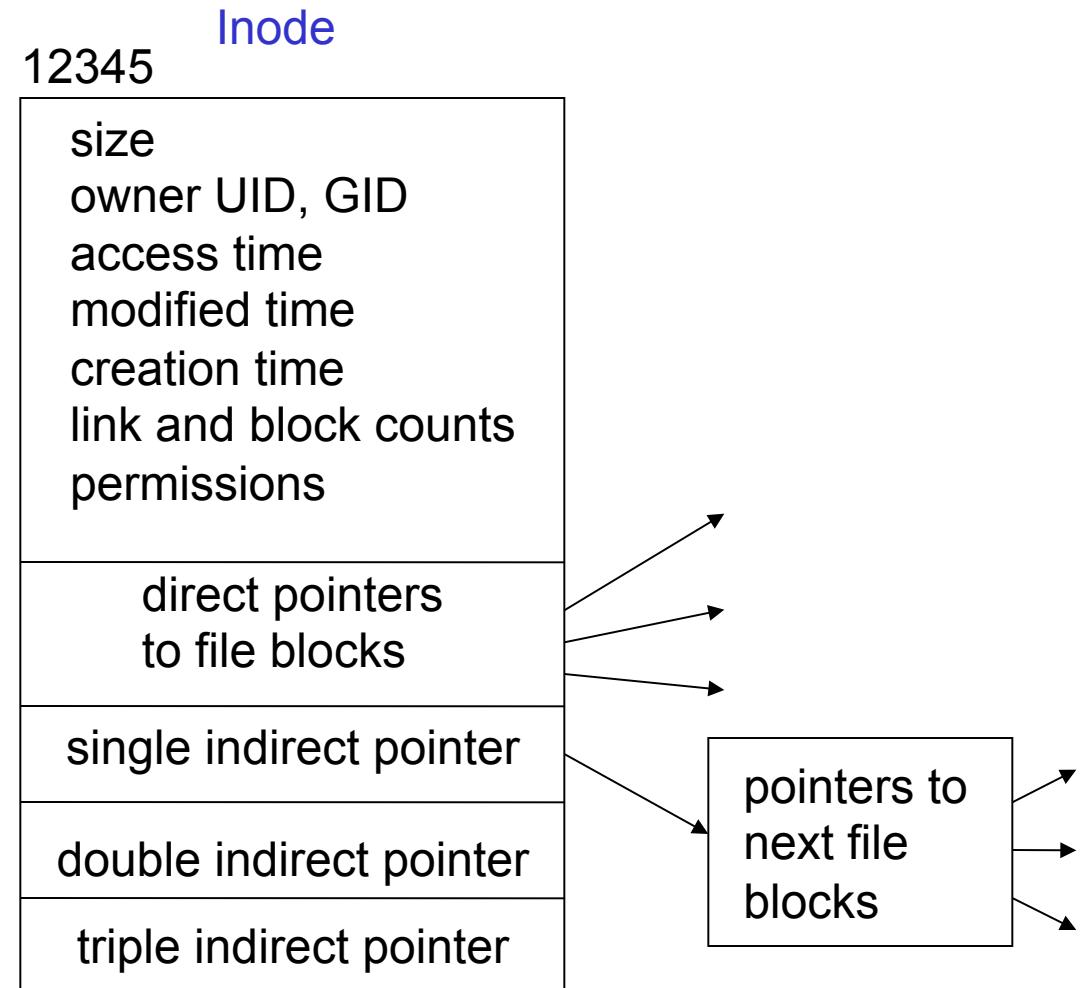


What is a file?

- A named collection of data with some attributes
 - Name
 - Owner (user and group)
 - Size
 - Permissions
 - Time of creation, last access, last modification
 - Location on disk
- How to get information on a file in Unix?
 - `ls -l`, `stat`

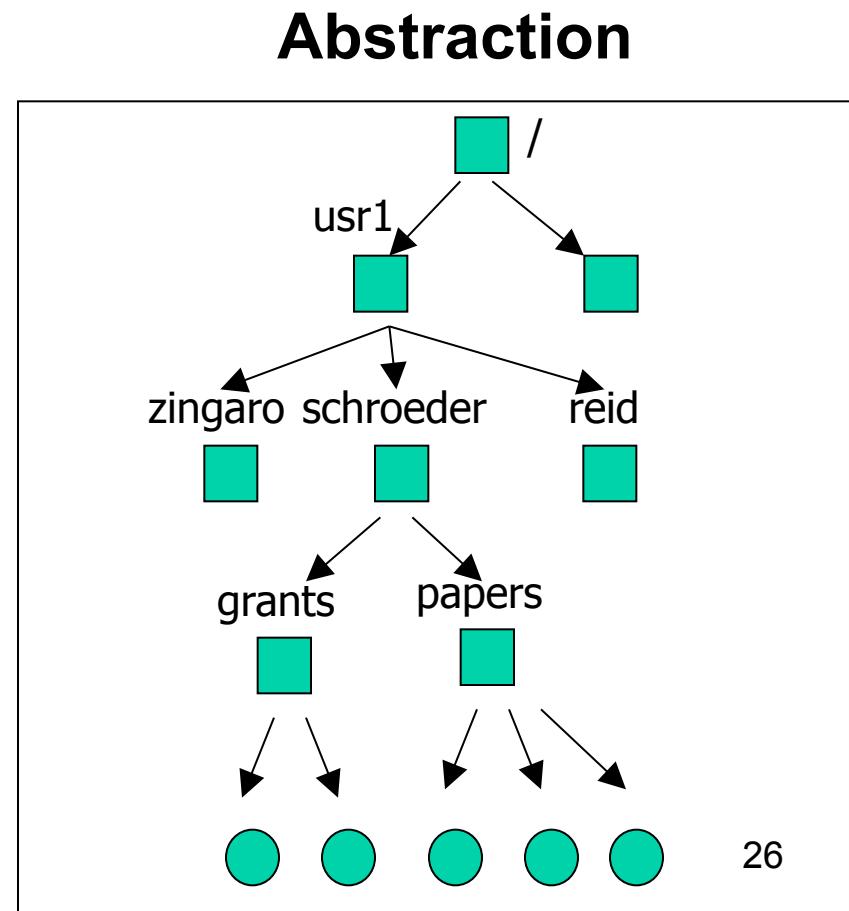
How does Unix implement a file?

- A data structure called **inode** stores the information about a file, including which disk blocks contain the file data.
- Identified by its **inode number**



What is a directory?

- A collection of files (and sub-directories)
- One special directory: the root directory (named “/”)
- In Unix: internally every directory is implemented as a file!



What is a directory?

- Internally, every directory is a file itself!
 - File with special content: *directory entries*
 - A directory entry maps a file name to an [inode](#).

Inode # 729482
(represents a directory)

size
owner UID, GID
access time
modified time
creation time
link and block counts
permissions
direct pointers to file blocks
single indirect pointer
double indirect pointer
triple indirect pointer

Inode# 12345
(represents file myfile)

Directory Entries	
Inode#	Filename
12345	myfile
12378	foofile
15384	barfile

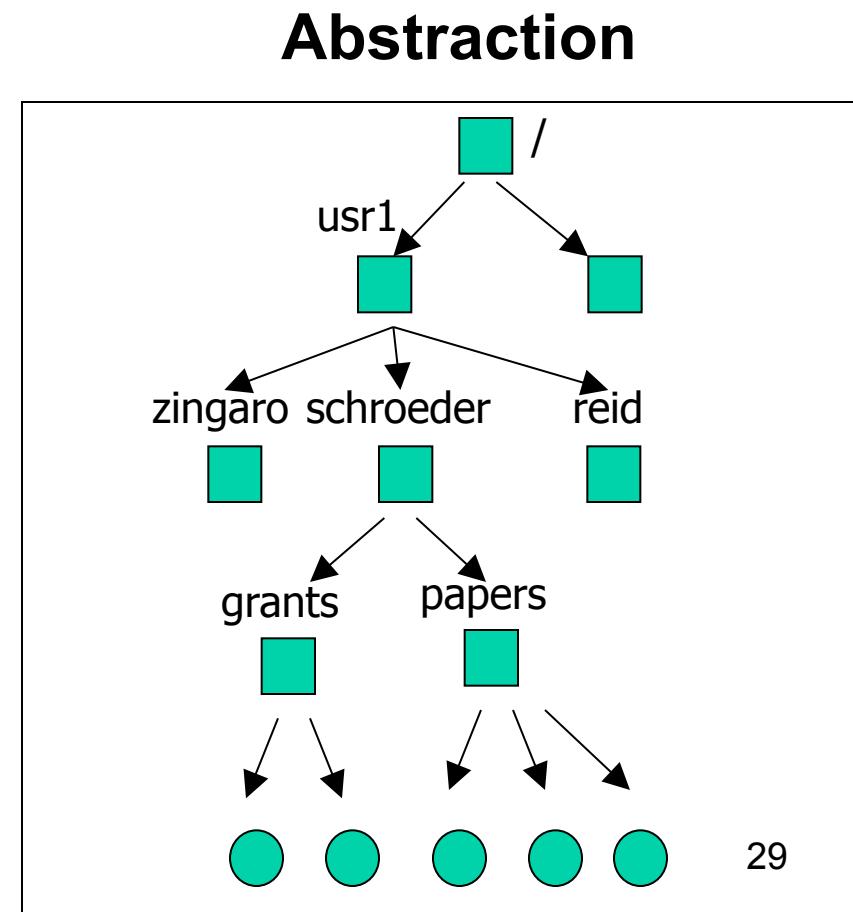
size
owner UID, GID
access time
modified time
creation time
link and block counts
permissions
direct pointers to file blocks
single indirect pointer
double indirect pointer
triple indirect pointer

Many more things are files ...

- “Everything is a file.”
- Unix provides a file interface for all Input/Output.
 - regular files
 - directories
 - devices
 - video (block)
 - keyboard (character)
 - sound (audio)
 - network (block)
- File interface = open, read, write, close

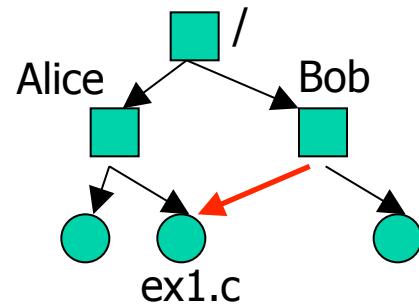
The directory hierarchy

- The directory tree is not really a tree as shown in the previous pictures...
 - But an acyclic graph
 - Why?



The directory hierarchy

- Sharing of files can be implemented by creating a new directory entry called a *link* : a pointer to another file or subdirectory
- An “ls” in Bob’s home directory will show ex1.c



Hard Links

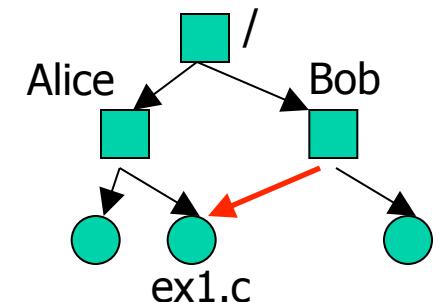
- Create with “`ln <target> <name of link>`”
- Second directory entry identical to the first
- What happens if anybody removes (`rm`) `ex1.c`?
 - A file is only removed if it no longer has any name/ hard links.

‘~Alice’ directory

File Name	Inode	Type
...
ex1.c	42	file
...

‘~Bob’ directory (hard link)

File Name	Inode	Type
...
ex1.c	42	file
...



Soft (symbolic) Links

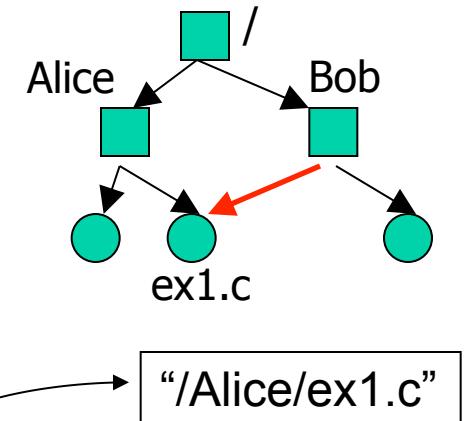
- Create with “`ln -s <target> <name of link>`”
- Directory entry points to small file containing the true path of the linked file
- What happens if anybody removes (`rm`) `ex1.c`?
 - If Alice removes `ex1.c`, dangling reference
 - If Bob removes `ex1.c`, only the entry in his directory is removed

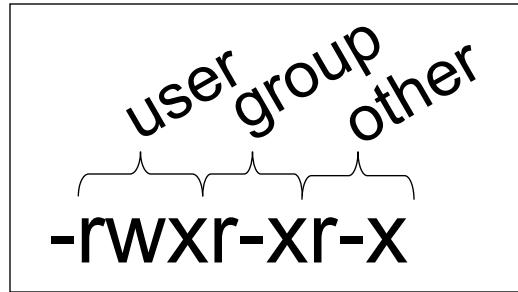
‘~Alice’ directory

File Name	Inode	Type
...
ex1.c	42	file
...

‘~Bob’ directory (soft link)

File Name	Inode	Type
...
ex1.c	529	link
...

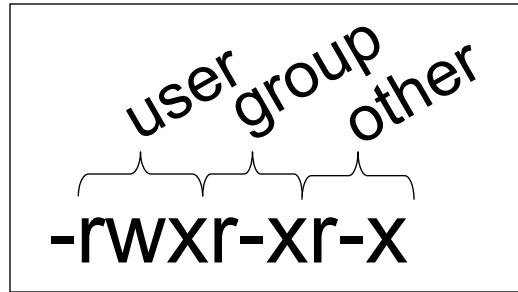




File Permissions

```
-rwxr-xr-x 1 schroel4  cmsusers  0 Jan  6 11:35 file1
-r--r--r-- 1 schroel4  cmsusers  0 Jan  6 11:35 file2
-rw------- 1 schroel4  cmsusers  0 Jan  6 11:35 file3
```

- Each file is owned by a particular user and group
- Each file has three permission entries
 - Permissions for owning user
 - Permissions for owning group
 - Permissions for all other users
- Each entries specify three permissions:
 - read, write, execute – pretty much what you think



Directory Permissions

```
drwxr-xr-x 2 schroel4    cmsusers 512 Jan  6 11:36 dir1/
dr-x---x---x 2 schroel4    cmsusers 512 Jan  6 11:36 dir2/
dr---r---r--- 2 schroel4    cmsusers 512 Jan  6 11:36 dir3/
```

- Directory permissions
 - read – you can run `ls` on the directory
 - write – you can create and delete files in the directory
 - execute – you can “pass through” the directory when searching subdirectories (you can `cd` into it) and access files (according to files’ permissions).

How can you change permissions?

- `chmod`
- Lots of ways, check `man chmod`
- E.g. give user owning file `fname` execute permissions:
`chmod u+x fname`
- Give all users in group read permissions
`chmod g+r fname`
- Remove read permissions from all users in group
`chmod g-r fname`
- Give all users all permissions
`chmod a+rwx fname`

That's it for today!