

CSCB09: Software Tools and Systems Programming

Bianca Schroeder

bianca@cs.toronto.edu

IC 460

The plan for today

- Processes
 - How to create new processes ...
- Why would you want to have a program that creates new processes?

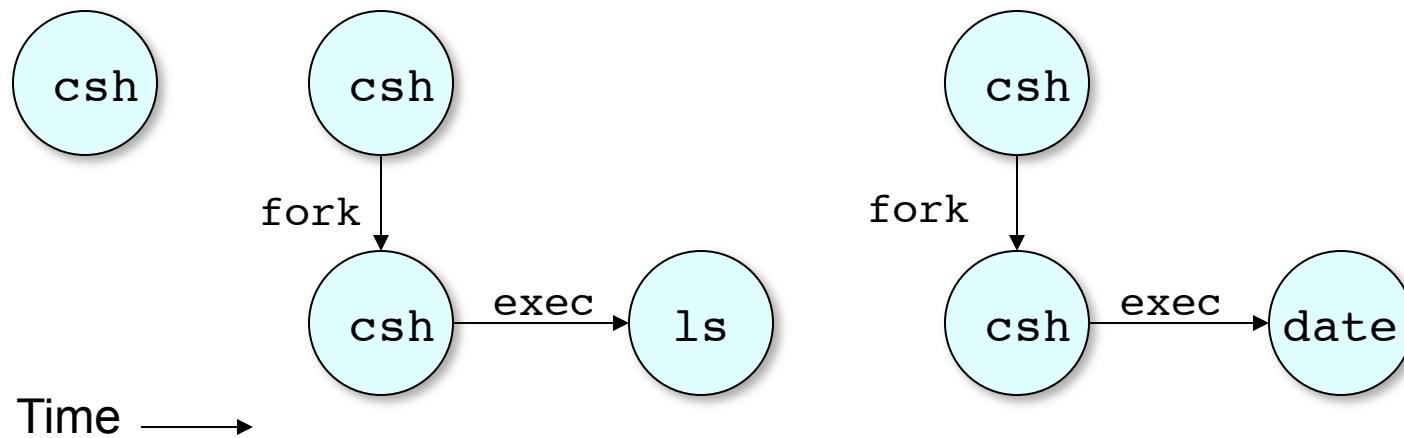
Creating new processes with fork

- Why would you want to have a program that creates new processes?
- Motivation: Hardware is not getting that much faster anymore
 - => use more hardware to solve a problem
 - Multiple processors/cores
 - Multiple machines
 - => Need to write a parallel program with multiple processes to exploit parallelism in hardware

Examples: Parallel programs

- Example: Web search engine
 - Web indices store counts of word occurrences for different files (web pages)
 - A search looks up a word in many different indices in parallel
- Example: A web server
 - Multiple processes serving web requests in parallel
- Example: A unix shell

Example: A shell

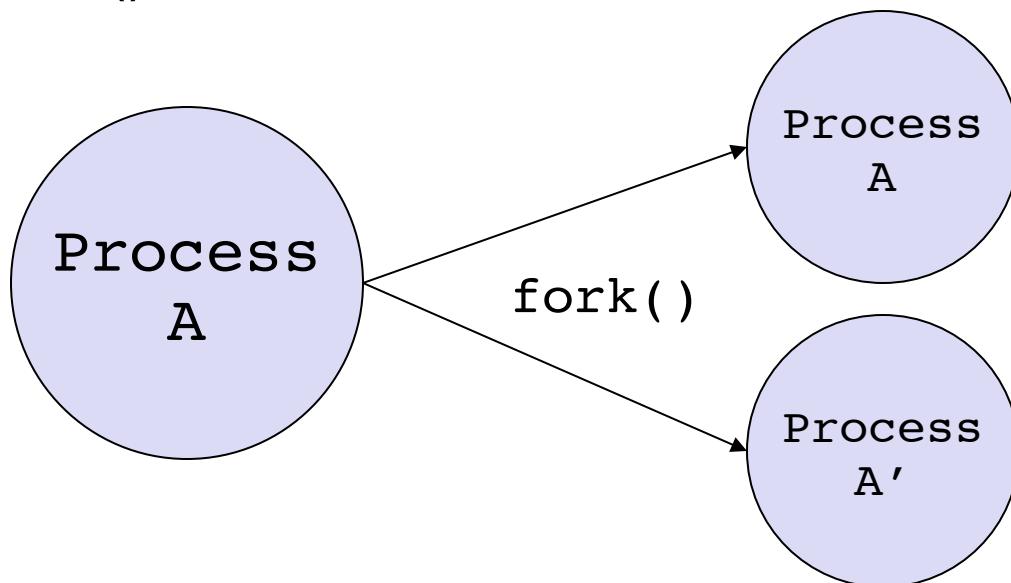


- When a command is typed, shell forks and then execs the typed command.

Processes can create other processes

- The fork system call creates a child process:

```
#include <unistd.h>
pid_t fork(void);
```
- Fork creates a **duplicate** of the currently running program.
- Both processes run concurrently and independently.
- After fork() both execute the next instruction after fork.



What is different between parent and forked child

- The child gets a new PID (process ID) and PPID
- The return value from the fork call is different:
 - On success
 - fork() returns 0 to the child
 - fork() returns the child's PID to the parent.
 - On failure (no child created) fork returns -1 to parent

Fork example

```
int main ()
{
    pid_t pid;
    pid = fork();
    if (pid < 0) {
        perror("fork()");
    } else if (pid > 0) {
        printf("parent\n");
    } else { /* pid == 0 */
        printf("child\n");
    }
    return 0;
}
```

Another fork example

Original process (parent)

```
int i; pid_t pid;  
i = 5;  
printf("%d\n", i);  
/* prints 5 */  
pid = fork();  
/* pid == 677 */  
if (pid > 0)  
    i = 6;  
printf("%d\n", i);
```

Child process

```
int i; pid_t pid;  
i = 5;  
printf("%d\n", i);  
  
pid = fork();  
/* pid == 0 */  
if (pid > 0)  
    i = 6;  
printf("%d\n", i);
```

- Each process gets their own copy of the address space
- After call to fork, variables are not shared.

Let's take a look at the first
work sheet!

When does a child terminate?

- Like any other program:
 - The program's main function returns
 - The program calls `exit`
 - `void exit(int status);`
 - (only the program calling `exit` will exit; not its children)
 - The program receives a signal that causes it to terminate
 - Will see next lecture what that means ...
- Programs have an exit status
 - Exit status is set by call to `exit` or `main`'s `return`
 - Exit status of most recent command stored in `$?` in most shells

wait()

- A parent might want to wait for a child to complete
- A parent might want to know the exit code of a child
- System call to wait for a child
 - `pid_t wait(int *status)`
- `wait` suspends execution of the calling process until one of its children terminates

wait()

- System call to wait for a child
 - `pid_t wait(int *status)`
- After calling `wait()` a process will:
 - block if all of its children are still running
 - return immediately with the PID of a terminated child, if there is a terminated child
 - return immediately with an error (-1) if it doesn't have any child processes.

Info returned by wait()

- System call to wait for a child
 - `pid_t wait(int *status)`
- Returns the pid of the terminated child or -1 on error
- `status` encodes the exit status of the child and how a child exited (normally or killed by signal)
- There are macros to process exit status:
 - `WIFEXITED` applied to `status` tells you if child terminated normally (i.e. by calling `exit` or `return`)
 - `WEXITSTATUS` gives you the exit status

Example using wait()

```
int main (void) {
    pid_t child;
    int status, exit_status;

    if ((child = fork()) == 0) {
        sleep (5);
        exit (8);
    }
    wait (&status);
    if (WIFEXITED(status)) {
        exit_status = WEXITSTATUS(status);
        printf ("Child %d done: %d\n", child, exit_status);
    }
    return 0;
}
```

treat as
an int

```
graph TD; A[treat as an int] --> B[pid_t fork()]; A --> C[pid_t wait(int *status)]; A --> D[void exit(int status)];
```

pid_t fork()
pid_t wait(int *status)
void exit(int status)

WIFEXITED(status)
WEXITSTATUS(status)