

# CSCB09: Software Tools and Systems Programming

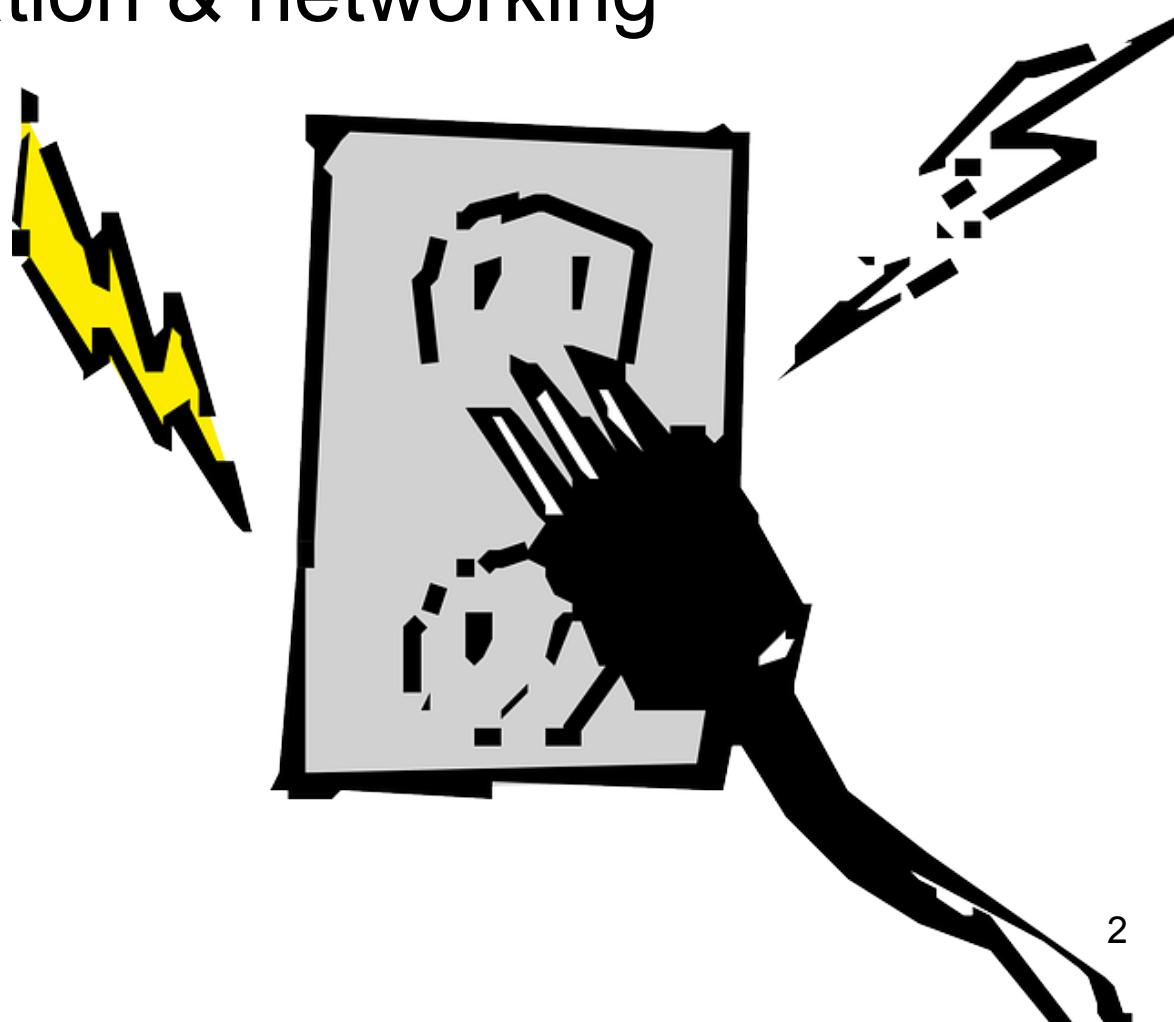
Bianca Schroeder

[bianca@cs.toronto.edu](mailto:bianca@cs.toronto.edu)

IC 460

# Today

- Communication & networking
- Sockets

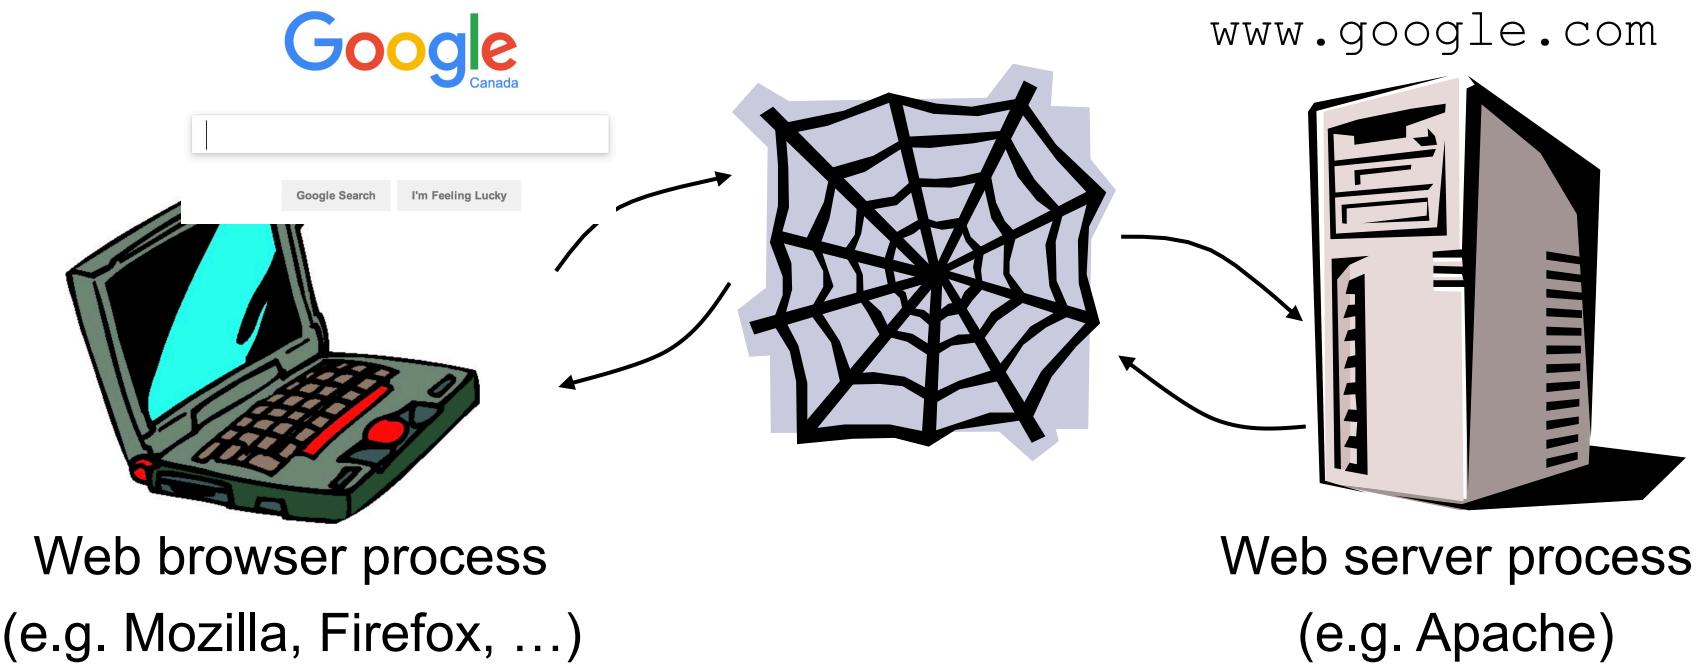


# Why?



Everything is networked now!

# Simple Web Request



- Two processes communicating: the web browser and the web server
- How exactly does this work?

# Client Server Model



**Server:** Always available at a known location



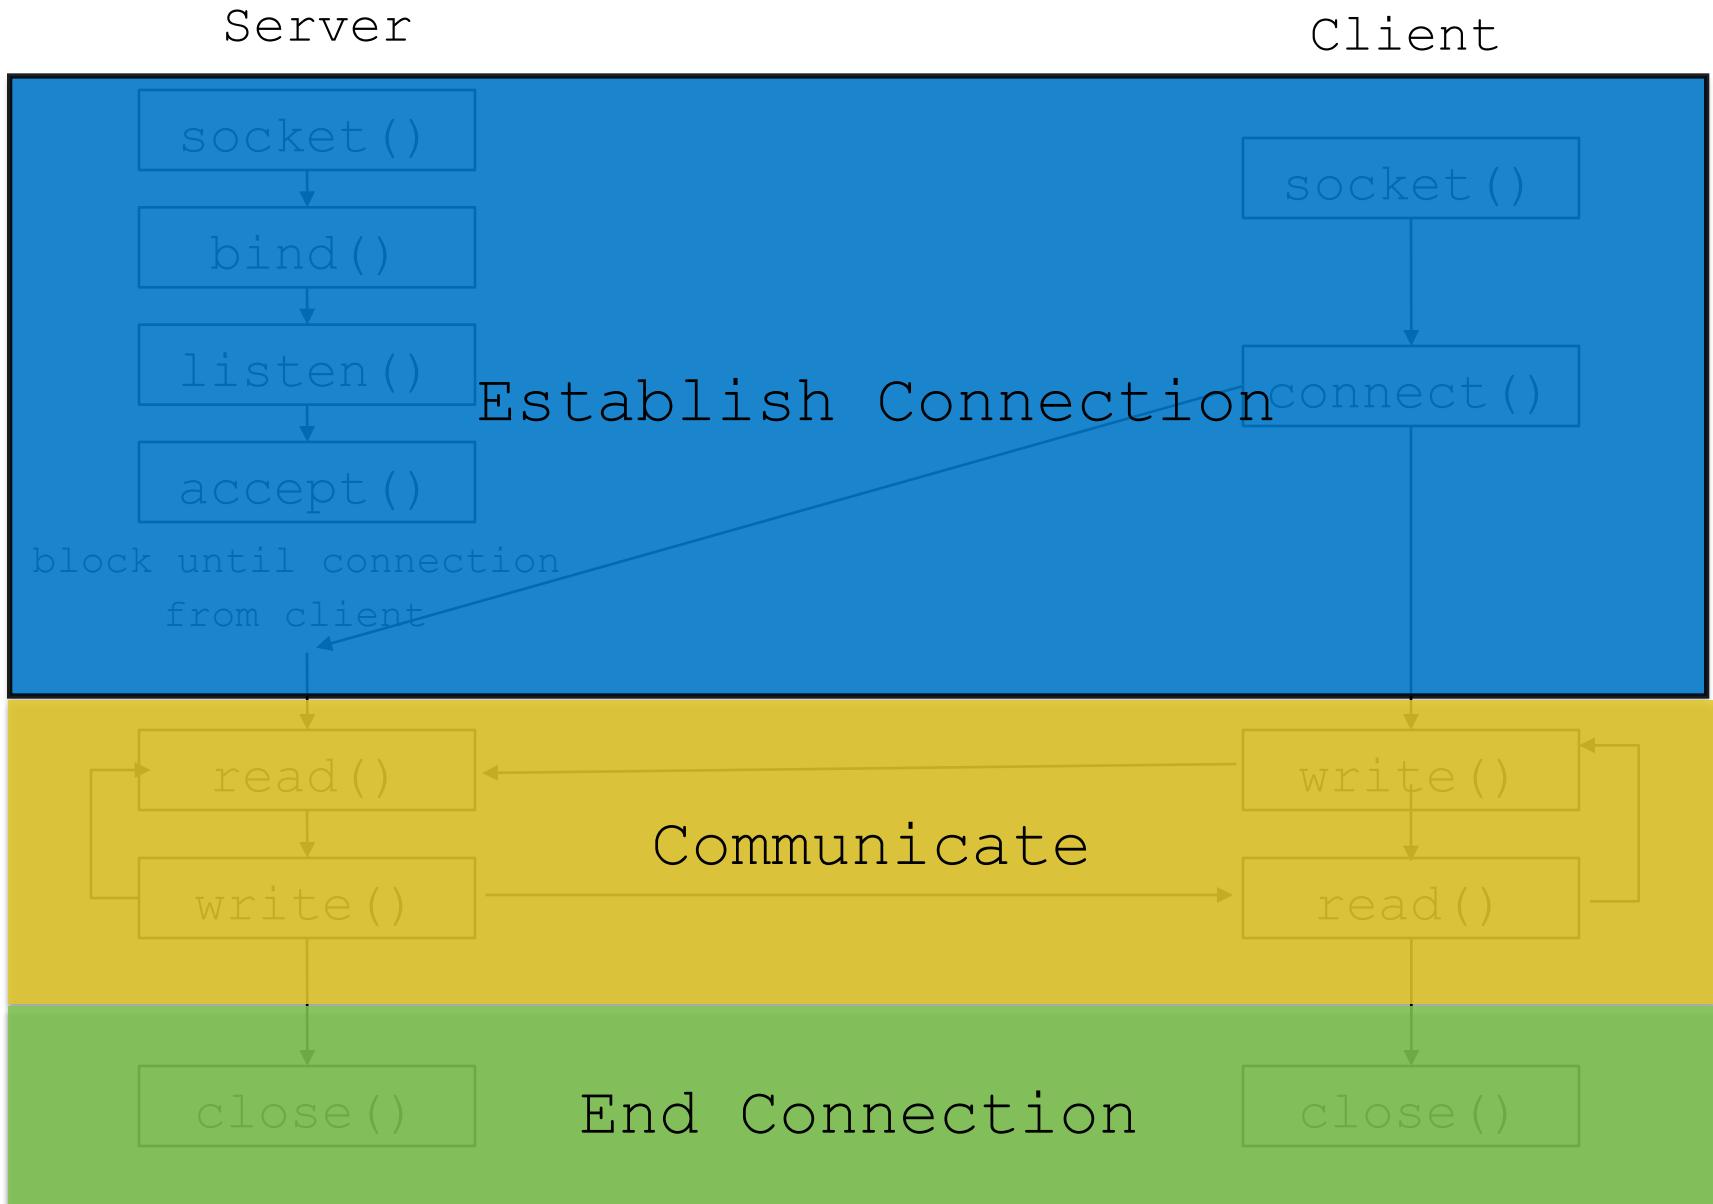
**Clients:** Show up and initiate the connection

# Some basic questions

- How do the server and browser talk to each other?
  - Pipes? Only for processes on same machine.
  - Solution: Sockets

# Sockets

- One form of communication between processes.
- Similar to pipes, except sockets can be used between processes on different machines and they are bi-directional.
  - Use file descriptors to refer to sockets.
  - Reading and writing similar to pipes
- The client and the server independently each create a socket using `socket()` system call

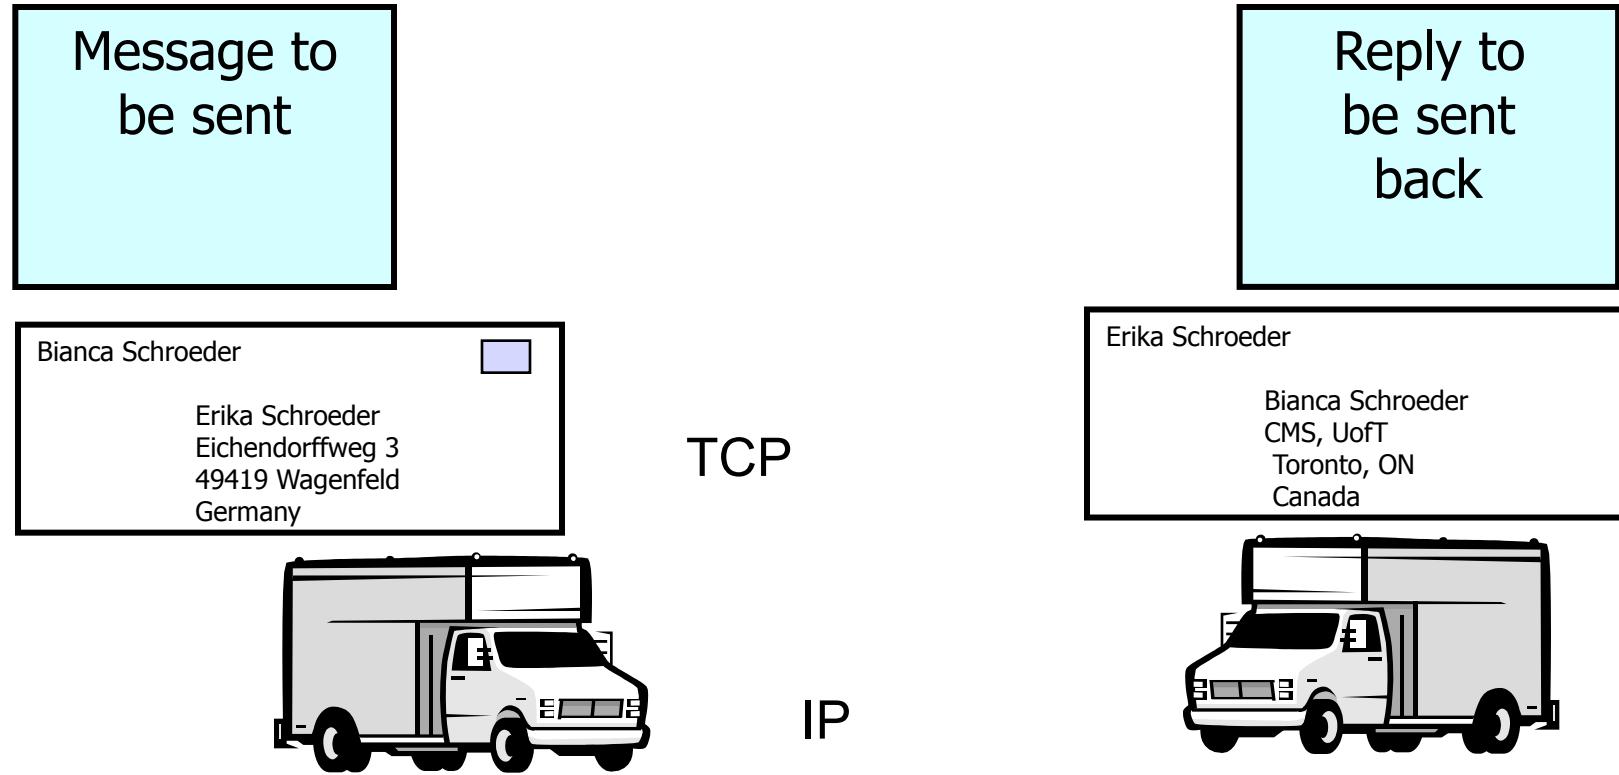


# The socket() system call

```
int socket(int family, int type,  
          int protocol);
```

- Returns a socket descriptor (similar to file descriptor)
- Arguments specify type of socket, including protocols to be used when processes exchange message on the socket
- Need to know a bit more about internet protocols ...

# Internet protocols



- The IP protocol enables routing of messages to the destination (based on correct address)
- The TCP protocol enables reliable, in-order delivery of messages. E.g. will resend lost message.

# The socket() system call

```
int socket(int family, int type,  
           int protocol);
```

- **family** specifies protocol family / domain:
  - `AF_INET` – IPv4
  - `AF_LOCAL` – Unix domain
- **type**
  - `SOCK_STREAM` (for reliable, in-order connection-based communication)
  - `SOCK_DGRAM`
  - `SOCK_RAW`
- **protocol**
  - set to `0` except for RAW sockets
- returns a socket descriptor

Let's set up a socket  
(Task 1 of today's worksheet) ...

# How do we specify the server?

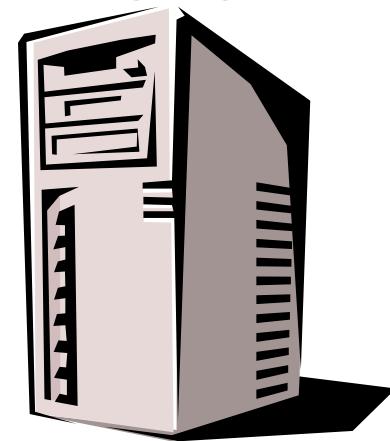


socket ()

Next: Client  
needs to  
specify which  
server to  
connect to

www.google.com

socket ()



IP address: 172.217.1.46

Port: 80

- Every computer on the Internet has an Internet address.
- Called an IP address (Internet Protocol)
- An IPv4 address is 4 numbers separated by dots.
- In addition client needs to specify, which process at server it wants to talk to: port number

# Addresses and Ports

- A **socket pair** is the two endpoints of the connection.
- An endpoint is identified by an **IP address and a port**.
- IPv4 addresses are 4 8-bit numbers:
  - 142.1.96.164 = mathlab.utsc.utoronto.ca
  - 142.1.96.30 = www.utsc.utoronto.ca
- Ports
  - because multiple processes can communicate with a single machine we need another identifier.

# More on Ports

- Well-known ports: 0-1023
  - 80 = http
  - 21 = ftp
  - 22 = ssh
  - 25 = smtp (mail)
- Registered ports: 1024-49151
  - 2709 = supermon
  - 26000 = quake
  - 3724 = world of warcraft
- Dynamic (private) ports: 49152-65535

# The connect() system call

```
int connect(int sockfd,  
           const struct sockaddr *servaddr,  
           socklen_t addrlen);
```

- Returns 0 on success.
- `servaddr` specifies which server to connect to.

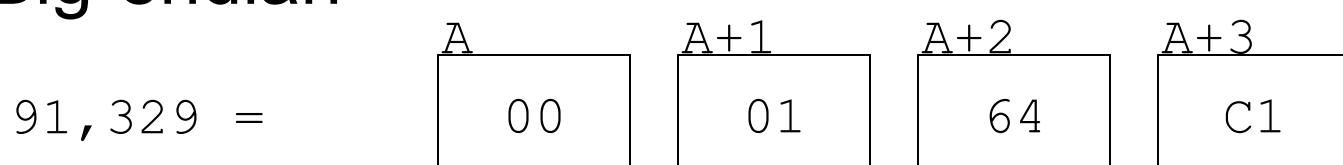
```
struct sockaddr_in {  
    short          sin_family; /*AF_INET */  
    u_short        sin_port;  
    struct in_addr sin_addr;  
    char           sin_zero[8];  
};
```

Let's see whether we can  
connect a socket  
(Task 2+3 of today's worksheet)

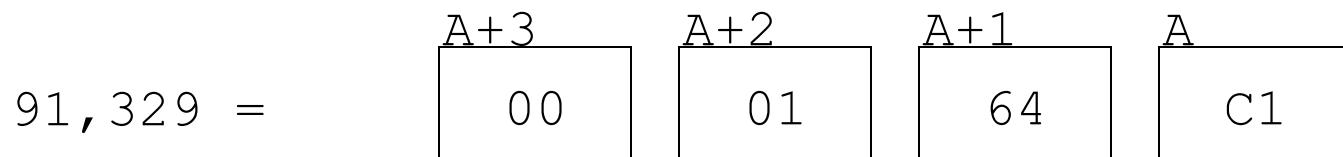
...

# Byte order

- Big-endian



- Little-endian



- Different hardware architectures have different endian-ness
- Intel is little-endian, and Sparc is big-endian

# Network byte order

- To communicate between machines with unknown or different “endian-ness” we convert numbers to network byte order (big-endian) before we send them.
- There are functions provided to do this:
  - `unsigned long htonl(unsigned long)`
  - `unsigned short htons(unsigned short)`
  - `unsigned long ntohl(unsigned long)`
  - `unsigned short ntohs(unsigned short)`

Now your socket should be ready  
for reading and writing  
(Task 4+5 of today's worksheet)

...