

# CSCB09: Software Tools and Systems Programming

Bianca Schroeder

[bianca@cs.toronto.edu](mailto:bianca@cs.toronto.edu)

IC 460

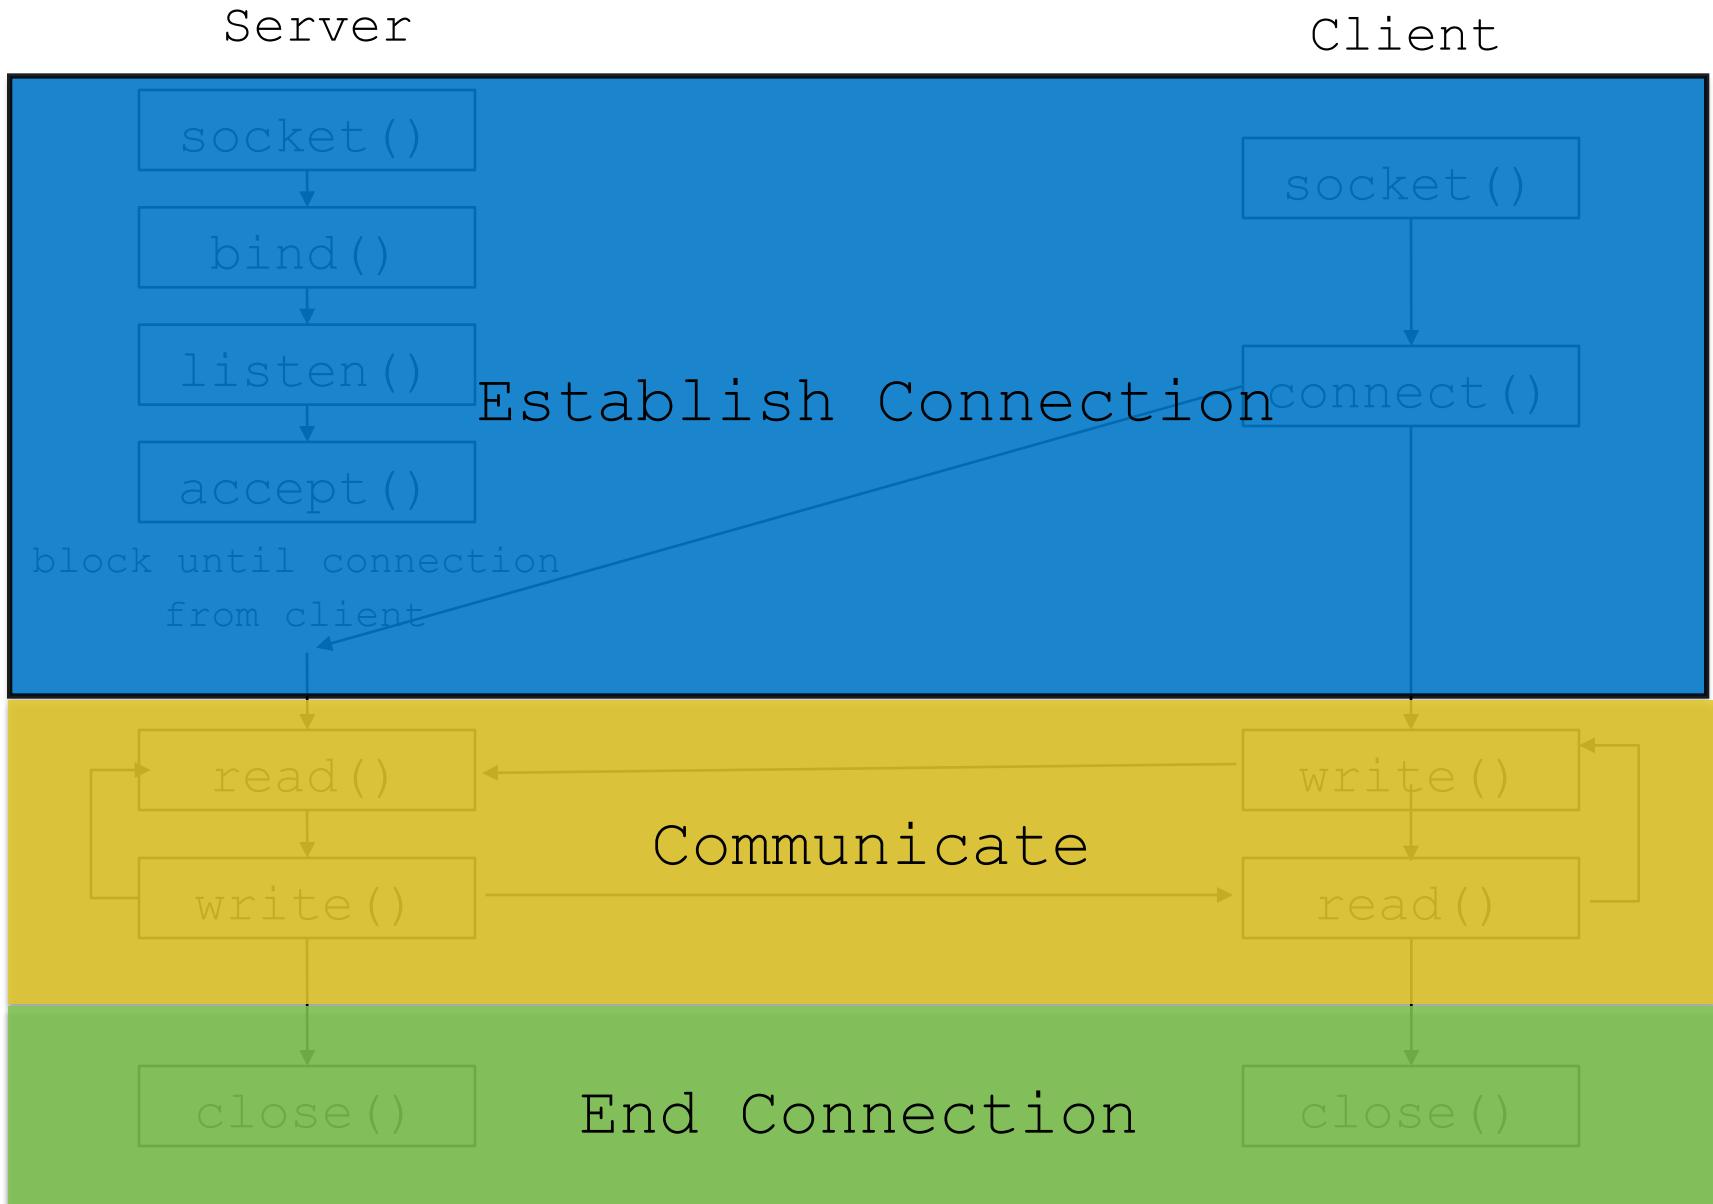
# COURSE EVALUATIONS

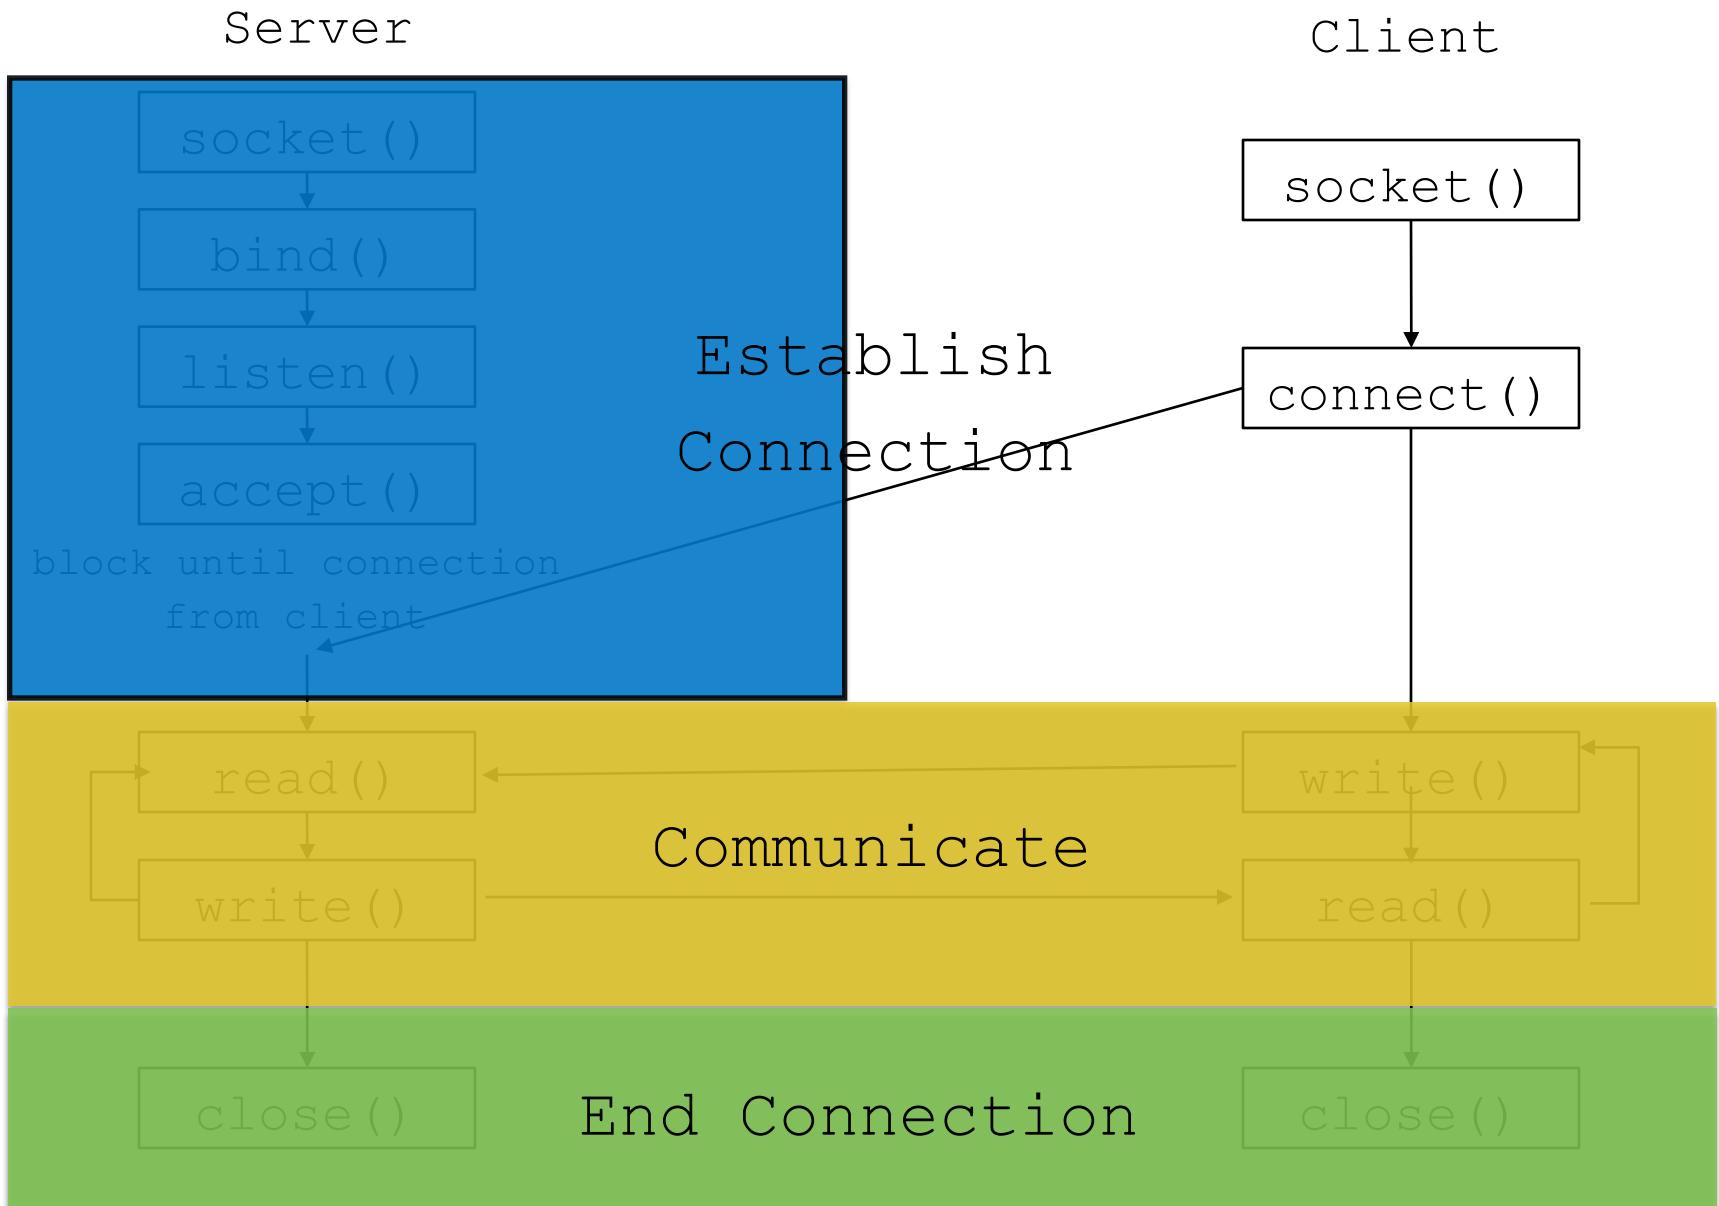
We care what  
you think,

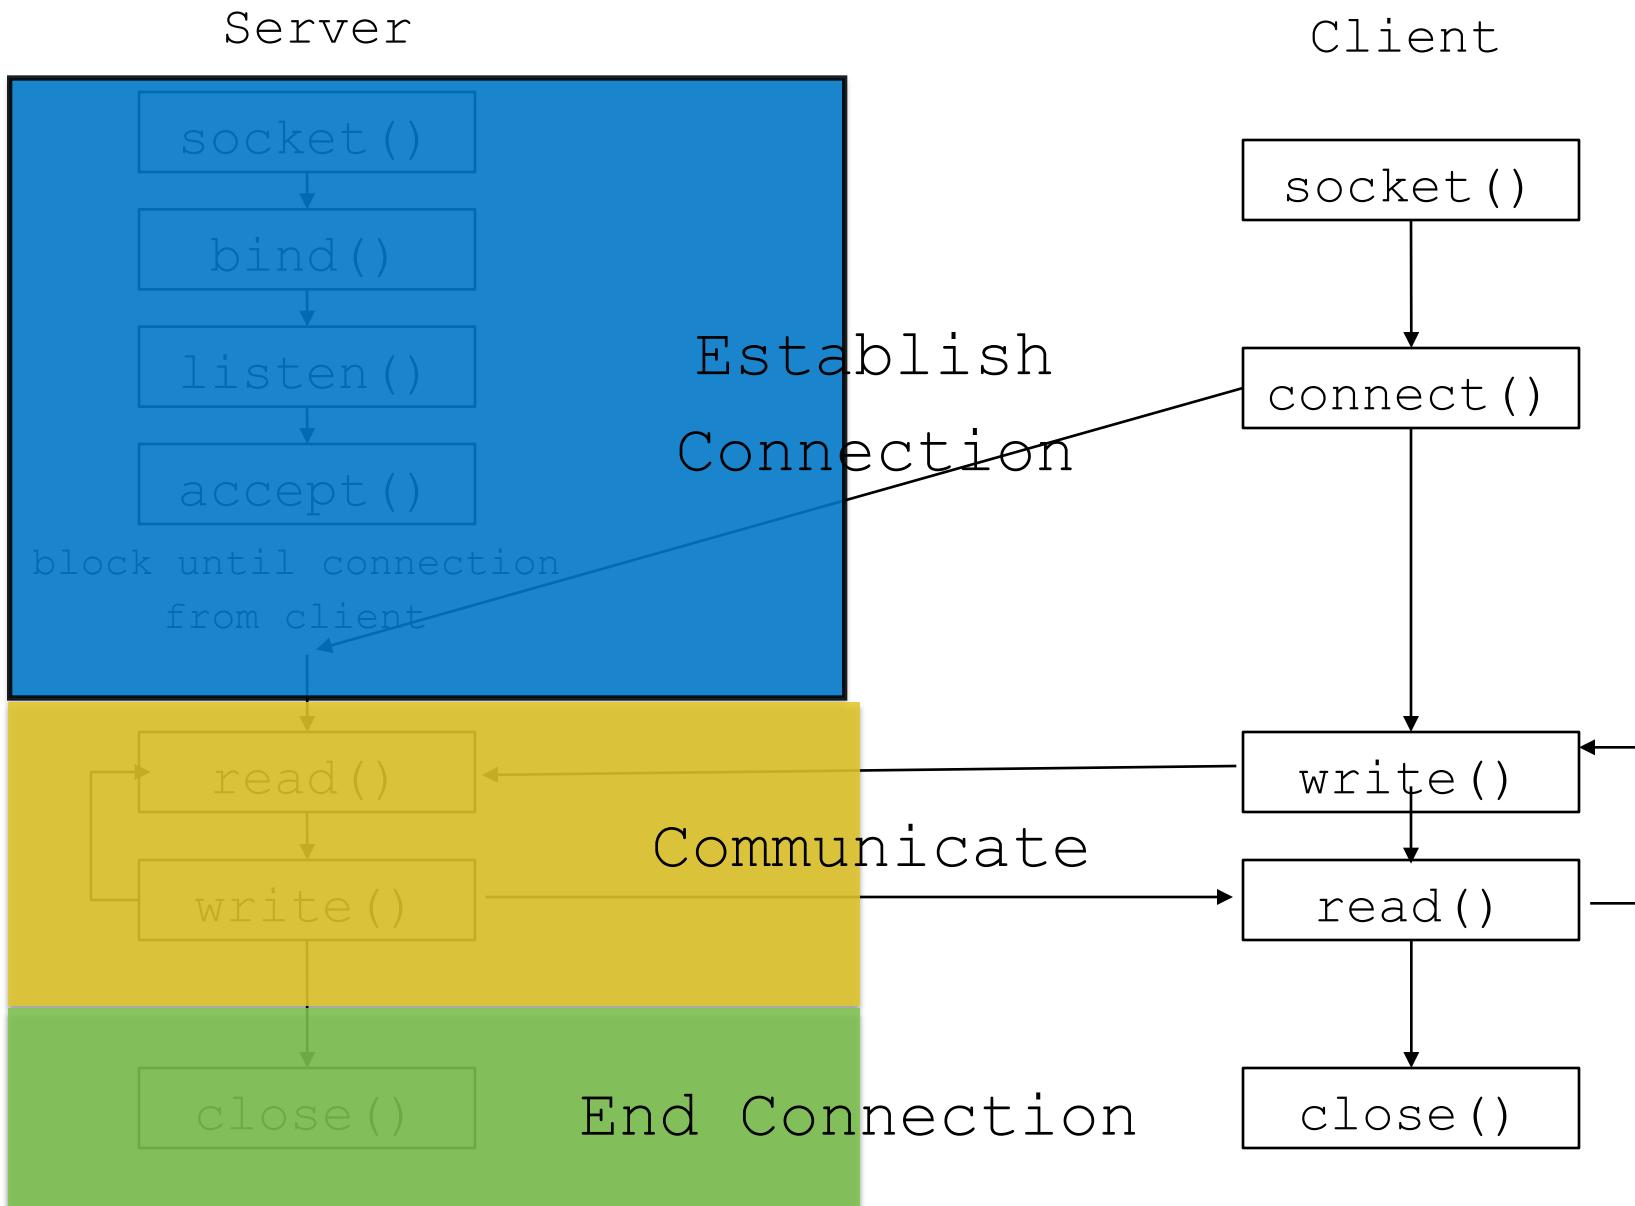


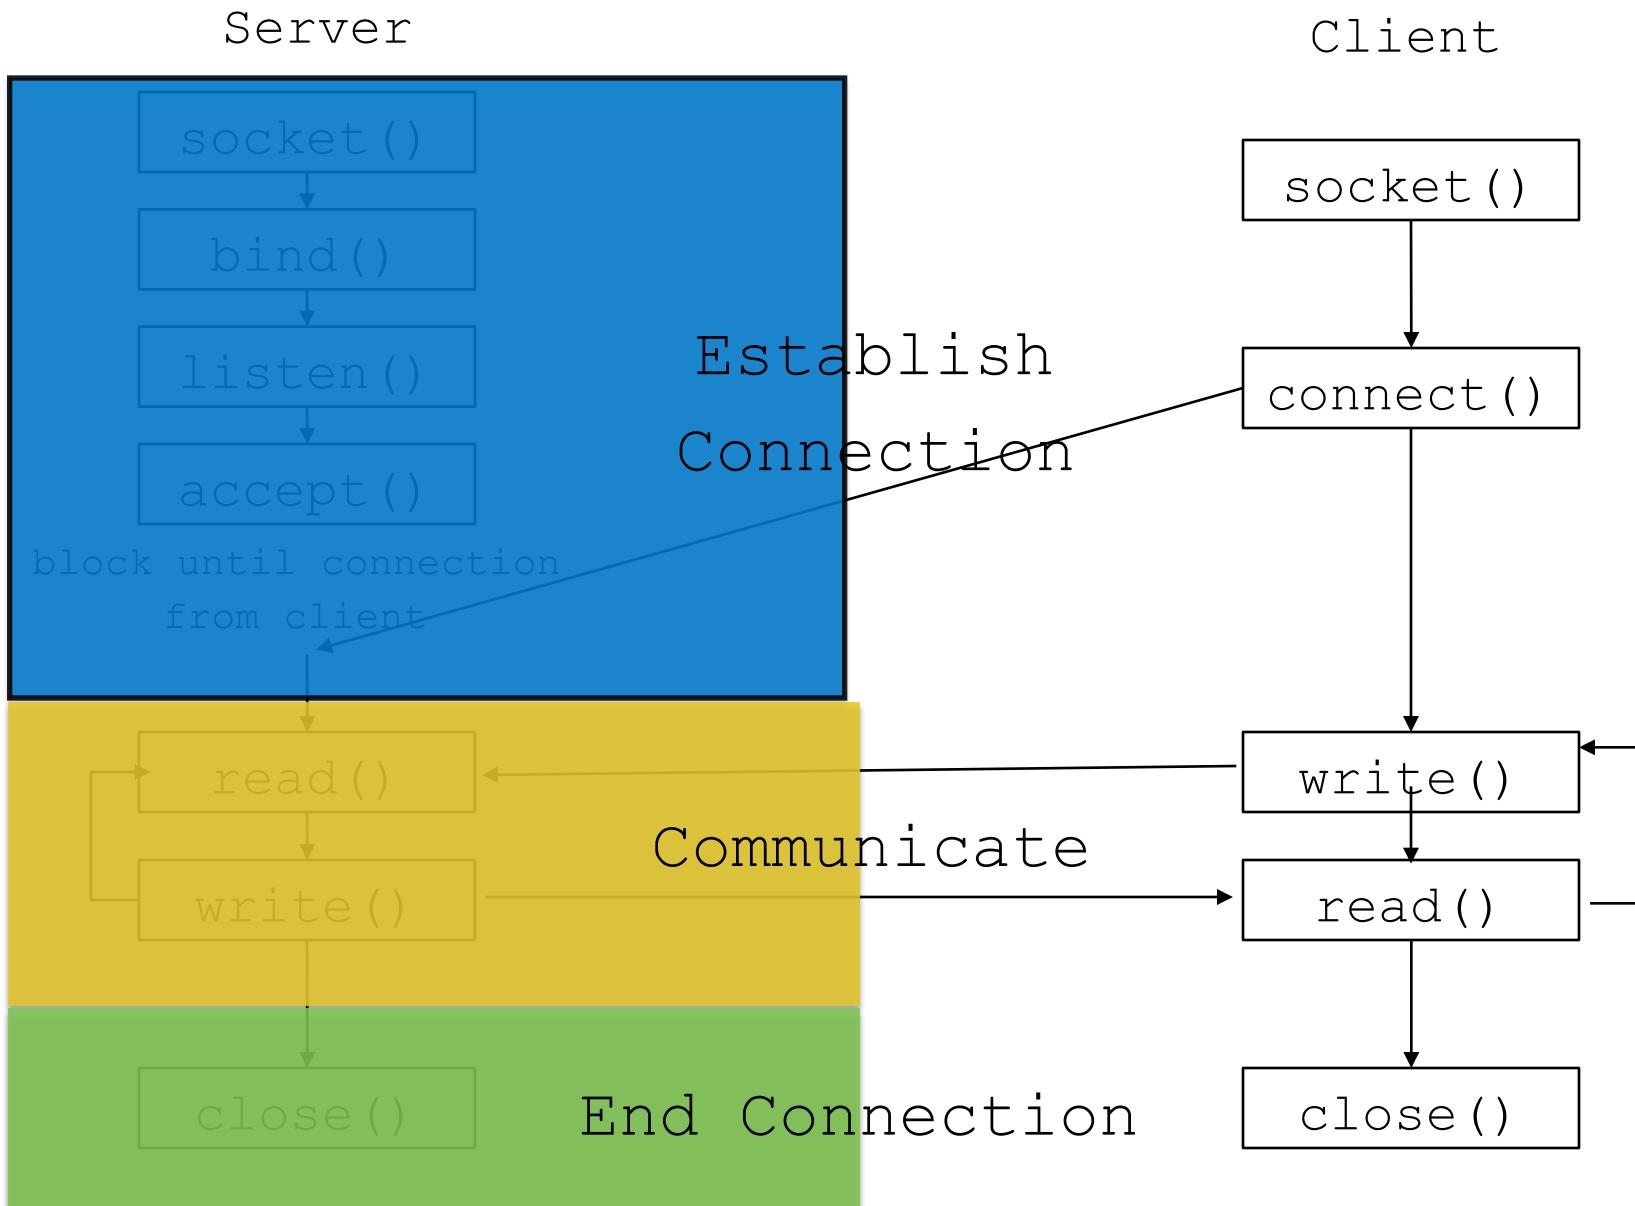
## TELL US!

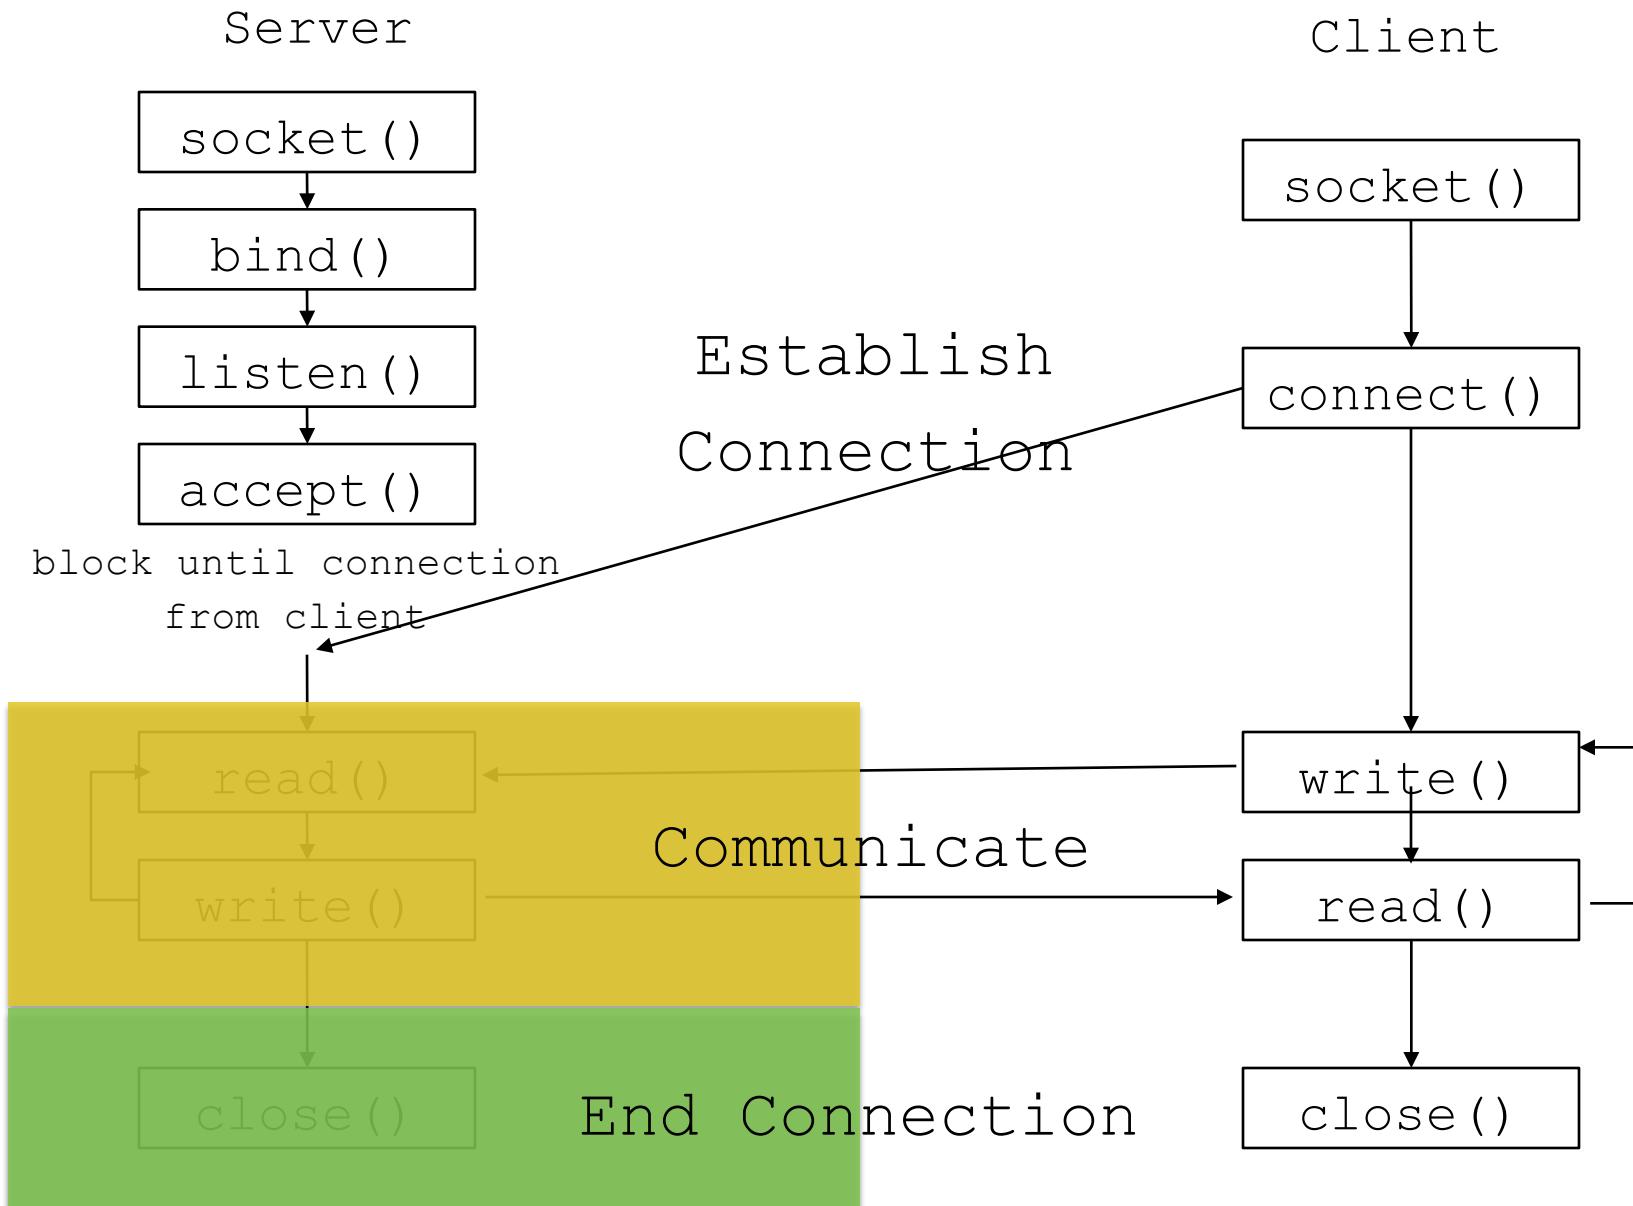


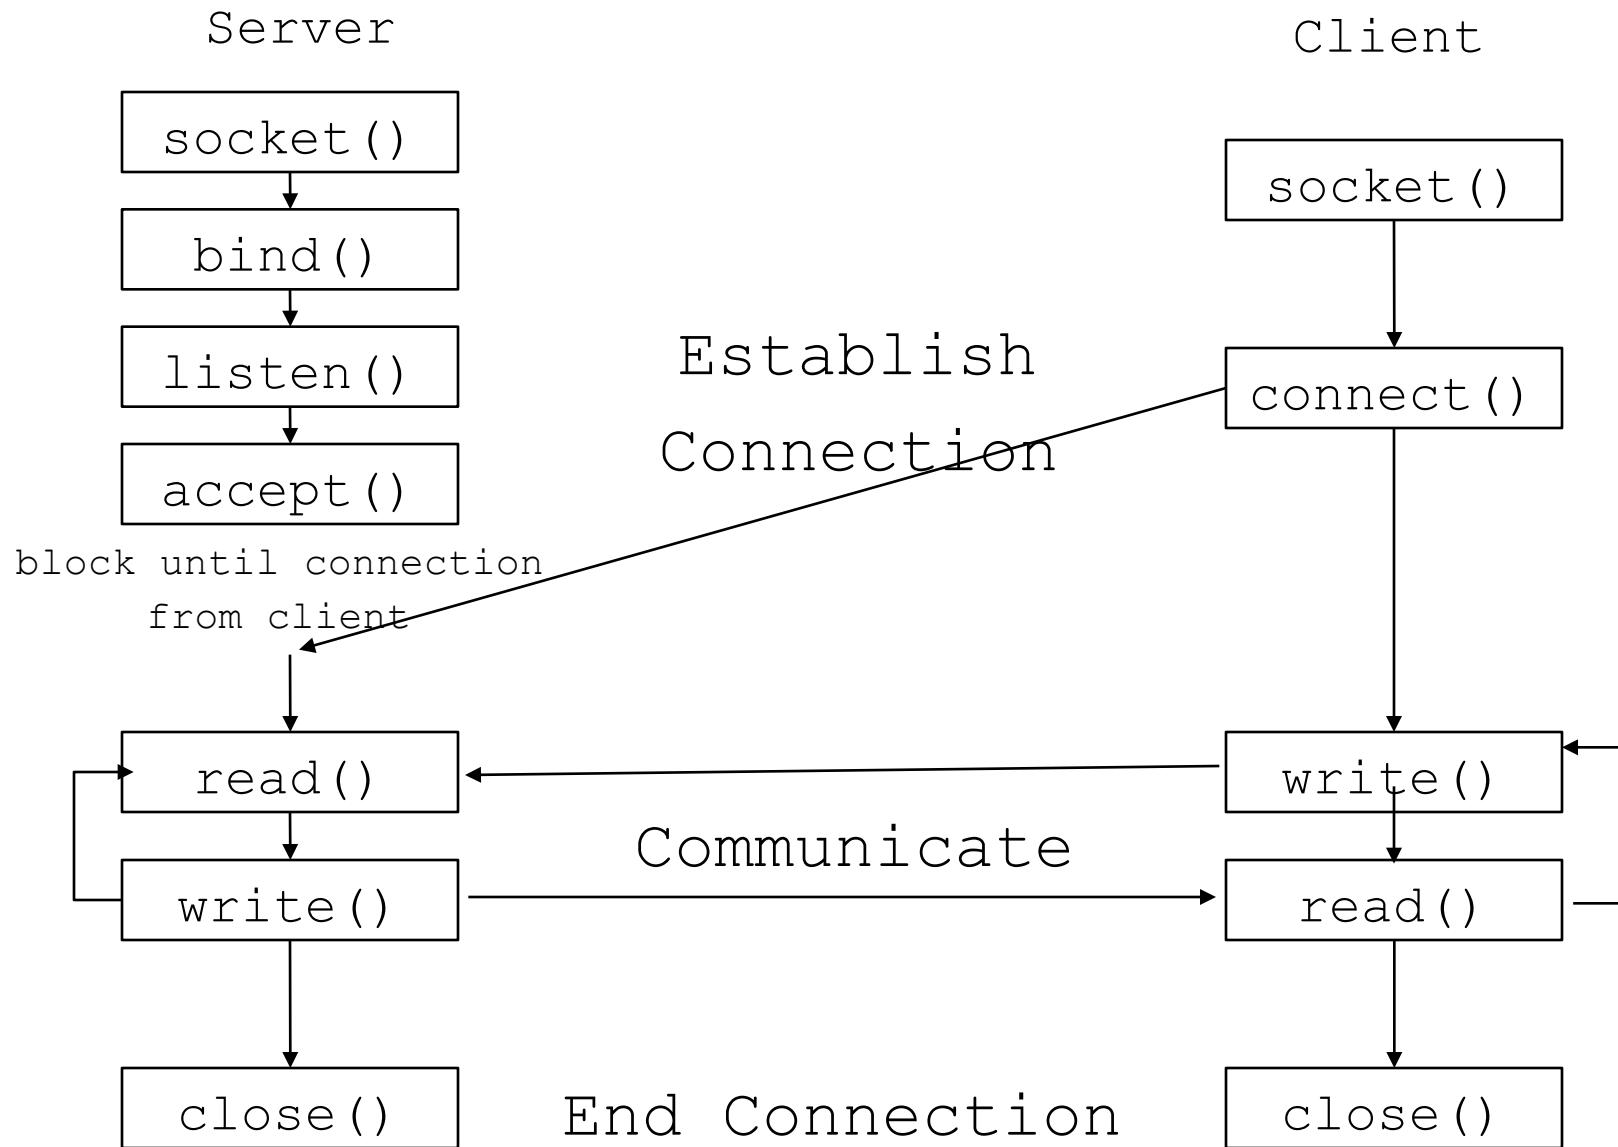












# Server side: socket() works same as for client

```
int socket(int family, int type,  
           int protocol);
```

- **family** specifies protocol family:
  - `AF_INET` – IPv4
  - `AF_LOCAL` – Unix domain
- **type**
  - `SOCK_STREAM`, `SOCK_DGRAM`, `SOCK_RAW`
- **protocol**
  - set to `0` except for RAW sockets
- returns a socket descriptor

# bind() socket to a name

```
int bind(int sockfd,  
        const struct sockaddr *servaddr,  
        socklen_t addrlen);
```

- Servaddr sets IP address & port at which to accept requests
- Returns negative value on error, zero on success
- ```
struct sockaddr_in {  
    short          sin_family; /*AF_INET */  
    u_short        sin_port;  
    struct in_addr sin_addr;  
    char           sin_zero[8]; /*filling*/  
};
```
- `sin_addr` can be set to `INADDR_ANY` to communicate on any network interface.

## Set up queue in kernel

```
int listen(int sockfd, int backlog)
```

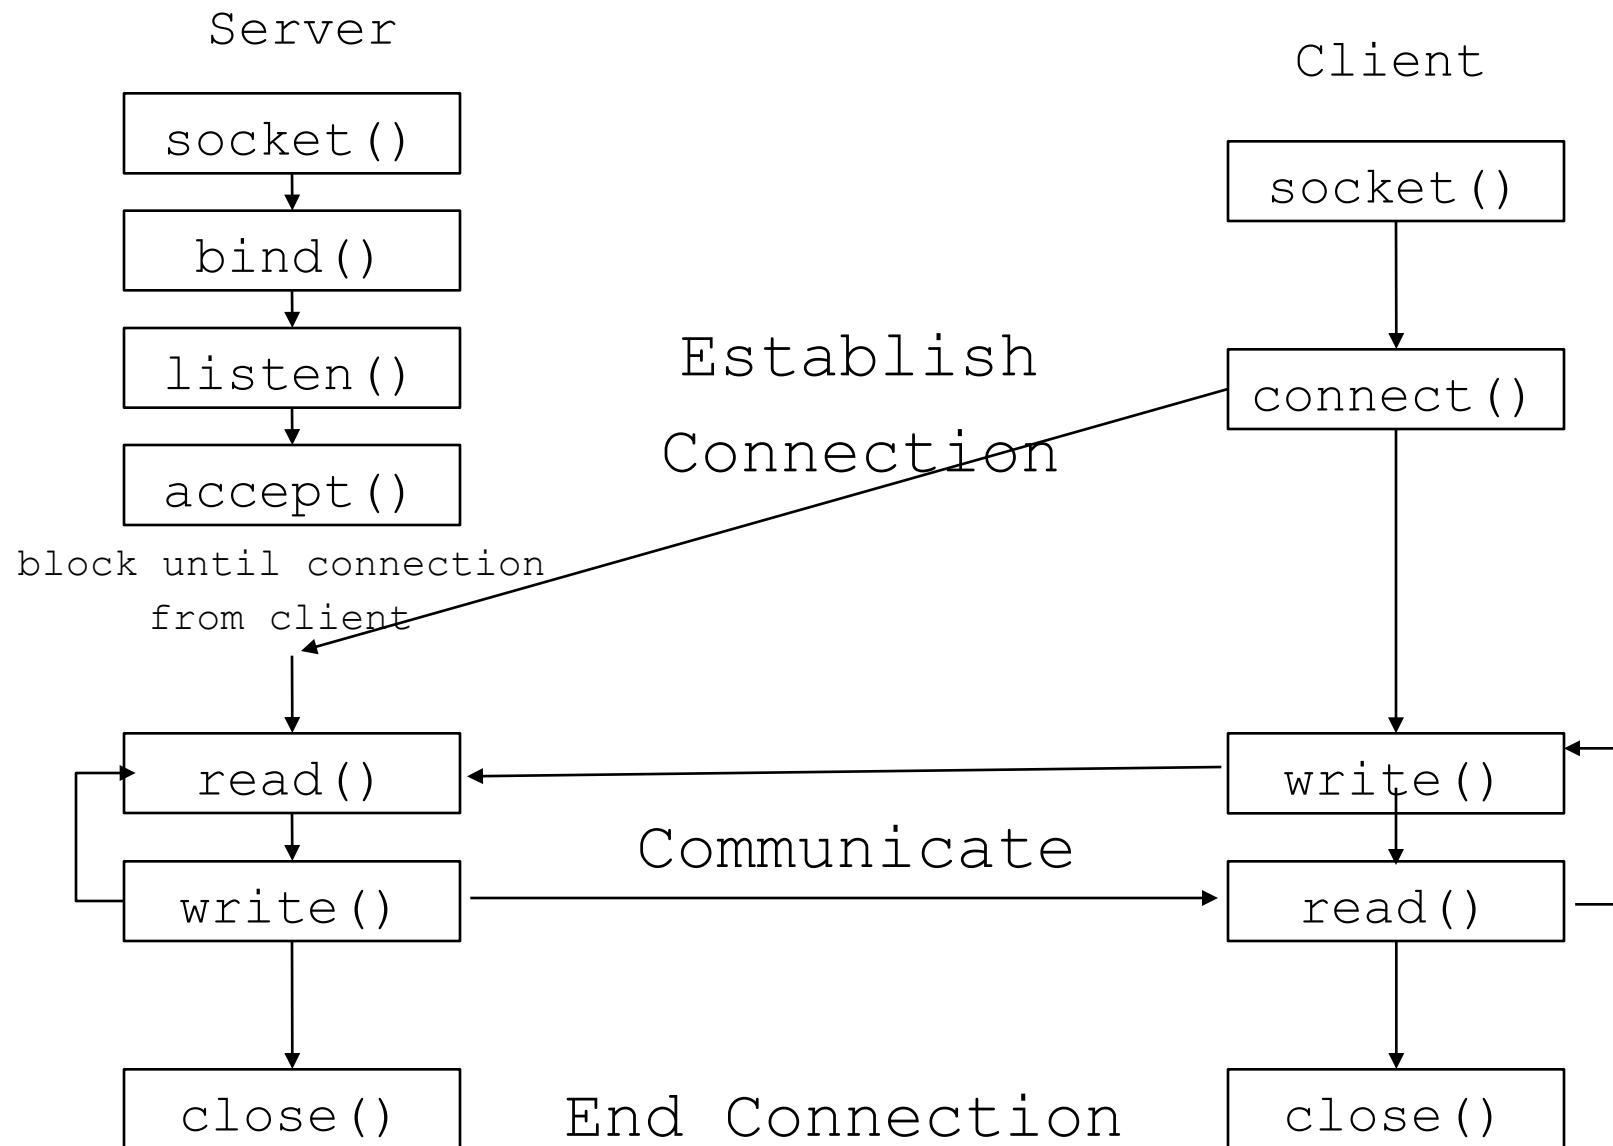
- after calling `listen`, a socket is ready to accept connections
- prepares a queue in the kernel where partially completed connections wait to be accepted.
- `backlog` is the maximum number of partially completed connections that the kernel should queue.

## Complete the connection

```
int accept(int sockfd,  
          struct sockaddr *cliaddr,  
          socklen_t *addrlen);
```

- blocks waiting for a connection (from the queue)
- returns a new descriptor which refers to the TCP connection with the client
  - This is the descriptor to be used for reads and writes from/to the client.
- **sockfd** is the listening socket
- **cliaddr** is used to return the address of the client

# Done!



# Next: Writing a server that can handle multiple clients

client1

server

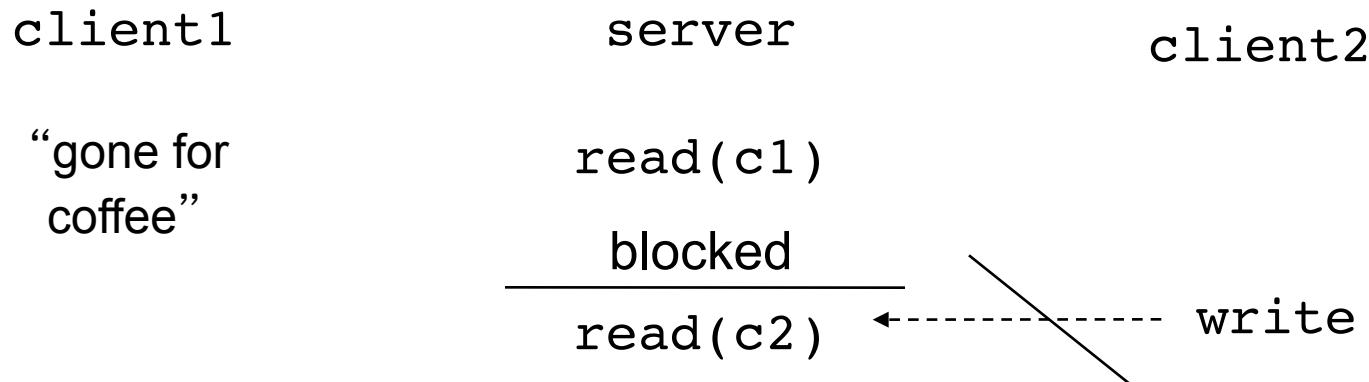
client2

```
while(1){  
    read(c1);  
    read(c2);  
}
```

The above is a draft for a server that reads from two clients.

Problems?

# The problem



When reading from multiple sources, blocking on one of the sources could be bad.

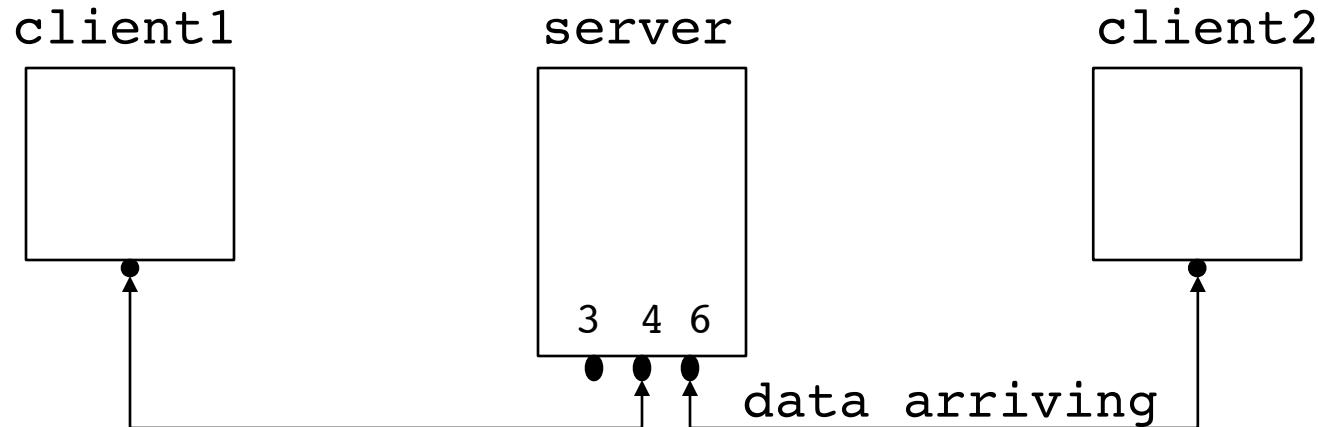
One solution: one process for every client. What are the pros and cons of this solution?

# The solution: Multiplexing

```
int select(int maxfdp1,  
          fd_set *readset,  
          fd_set *writeset,  
          fd_set *exceptset,  
          const struct timeval *timeout);
```

- A call to select returns when one of the file descriptors in one of the sets is ready for I/O.
- Upon return, the file descriptor set contains the file descriptors that are ready for I/O.
- If timeout is not NULL, then select returns when a descriptor is ready or timeout time has passed.
- If timeout is 0, select returns immediately after checking descriptors.

# Descriptor sets



|        | fd0           | fd1 | fd2 | fd3 | fd4 | fd5 | fd6 |
|--------|---------------|-----|-----|-----|-----|-----|-----|
| allset | 0             | 0   | 0   | 1   | 1   | 0   | 1   |
|        | maxfd + 1 = 7 |     |     |     |     |     |     |

After select (assuming client2 has written some data):

| rset | 0 | 0 | 0 | 0 | 0 | 0 | 1 |
|------|---|---|---|---|---|---|---|
|------|---|---|---|---|---|---|---|

# Readiness

## Ready to read when

- there is data in the receive buffer to be read

- end-of-file state on file descriptor

- the socket is a listening socket and there is a connection pending

- a socket error is pending

## Ready to write when

- there is space available in the write buffer

- a socket error is pending

## Exception condition pending when

- TCP out-of-band data

We are typically interested in when bytes are available to be read, but sometimes we use select on write or exception sets

# Descriptor sets

- Implementation is hidden in the `fd_set` data type
- `FD_SETSIZE` is the number of descriptors in the data type
- `maxfdp1` specifies the number of descriptors to test

Macros:

```
void FD_ZERO(fd_set *fdset);
void FD_SET(int fd, fd_set *fdset);
void FD_CLR(int fd, fd_set *fdset);
int FD_ISSET(int fd, fd_set *fdset);
```

# Let's put this into practice on the worksheet

```
FD_ZERO(fd_set *set)  
FD_SET(int fd, fd_set *set)  
FD_ISSET(int fd, fd_set *set)
```

```
select(int max_fd,  
       fd_set *read_fds,  
       fd_set *write_fds,  
       fd_set *except_fds,  
       struct timeval *timeout)
```

**STDIN\_FILENO**      }  
**STDOUT\_FILENO**    }

File descriptors for stdin and stdout